



Zero to Thousands TPS: Private Cloud Operations Workshop

Maudrit Martinez
Apigee

Paul Mibus
Apigee

Agenda

9:30 – 10:00

Apigee Overview

10:00 – 11:00

Apigee Private Cloud Architecture

BREAK (15 MINS)

11:15 – 12:00

Apigee Private Cloud Architecture (cont)

BREAK – LUNCH (60 MINS)

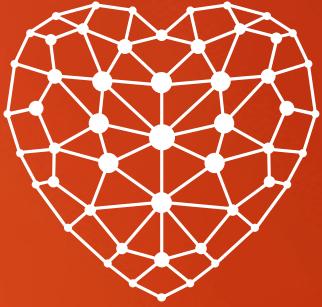
13:00 – 15:45

Apigee Edge Install

BREAK (15 MINS)

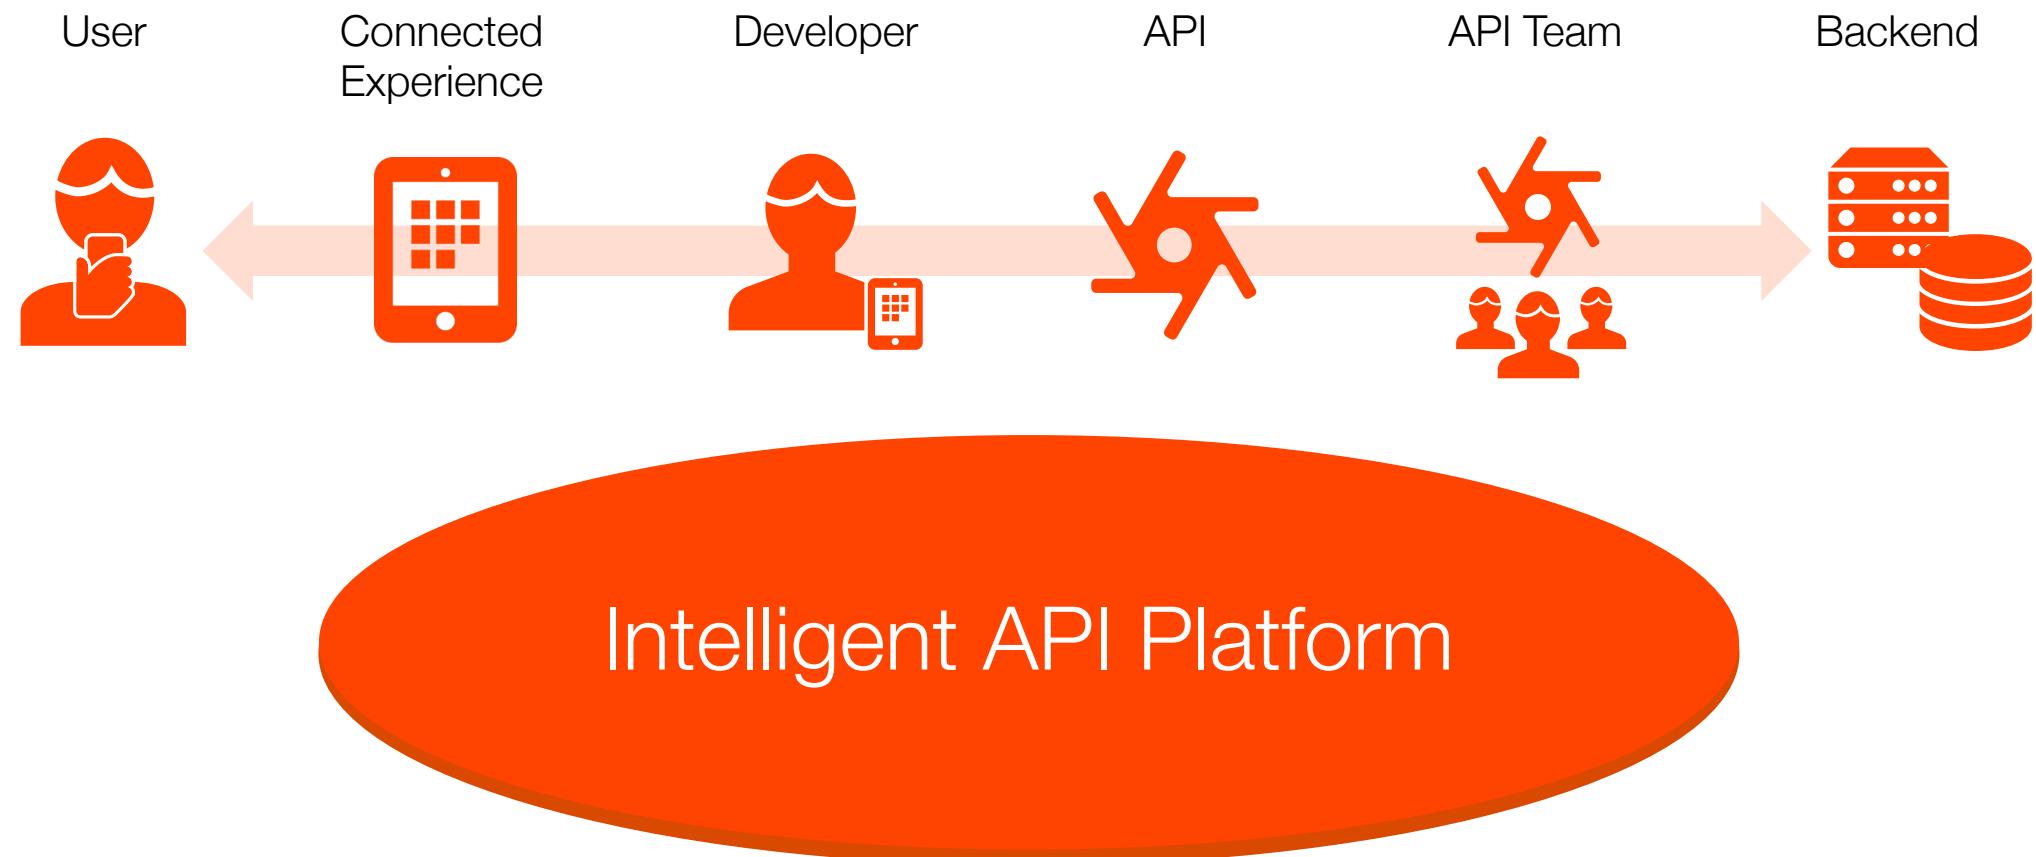
16:00 – 17:30

Platform Operations



Apigee Edge Overview

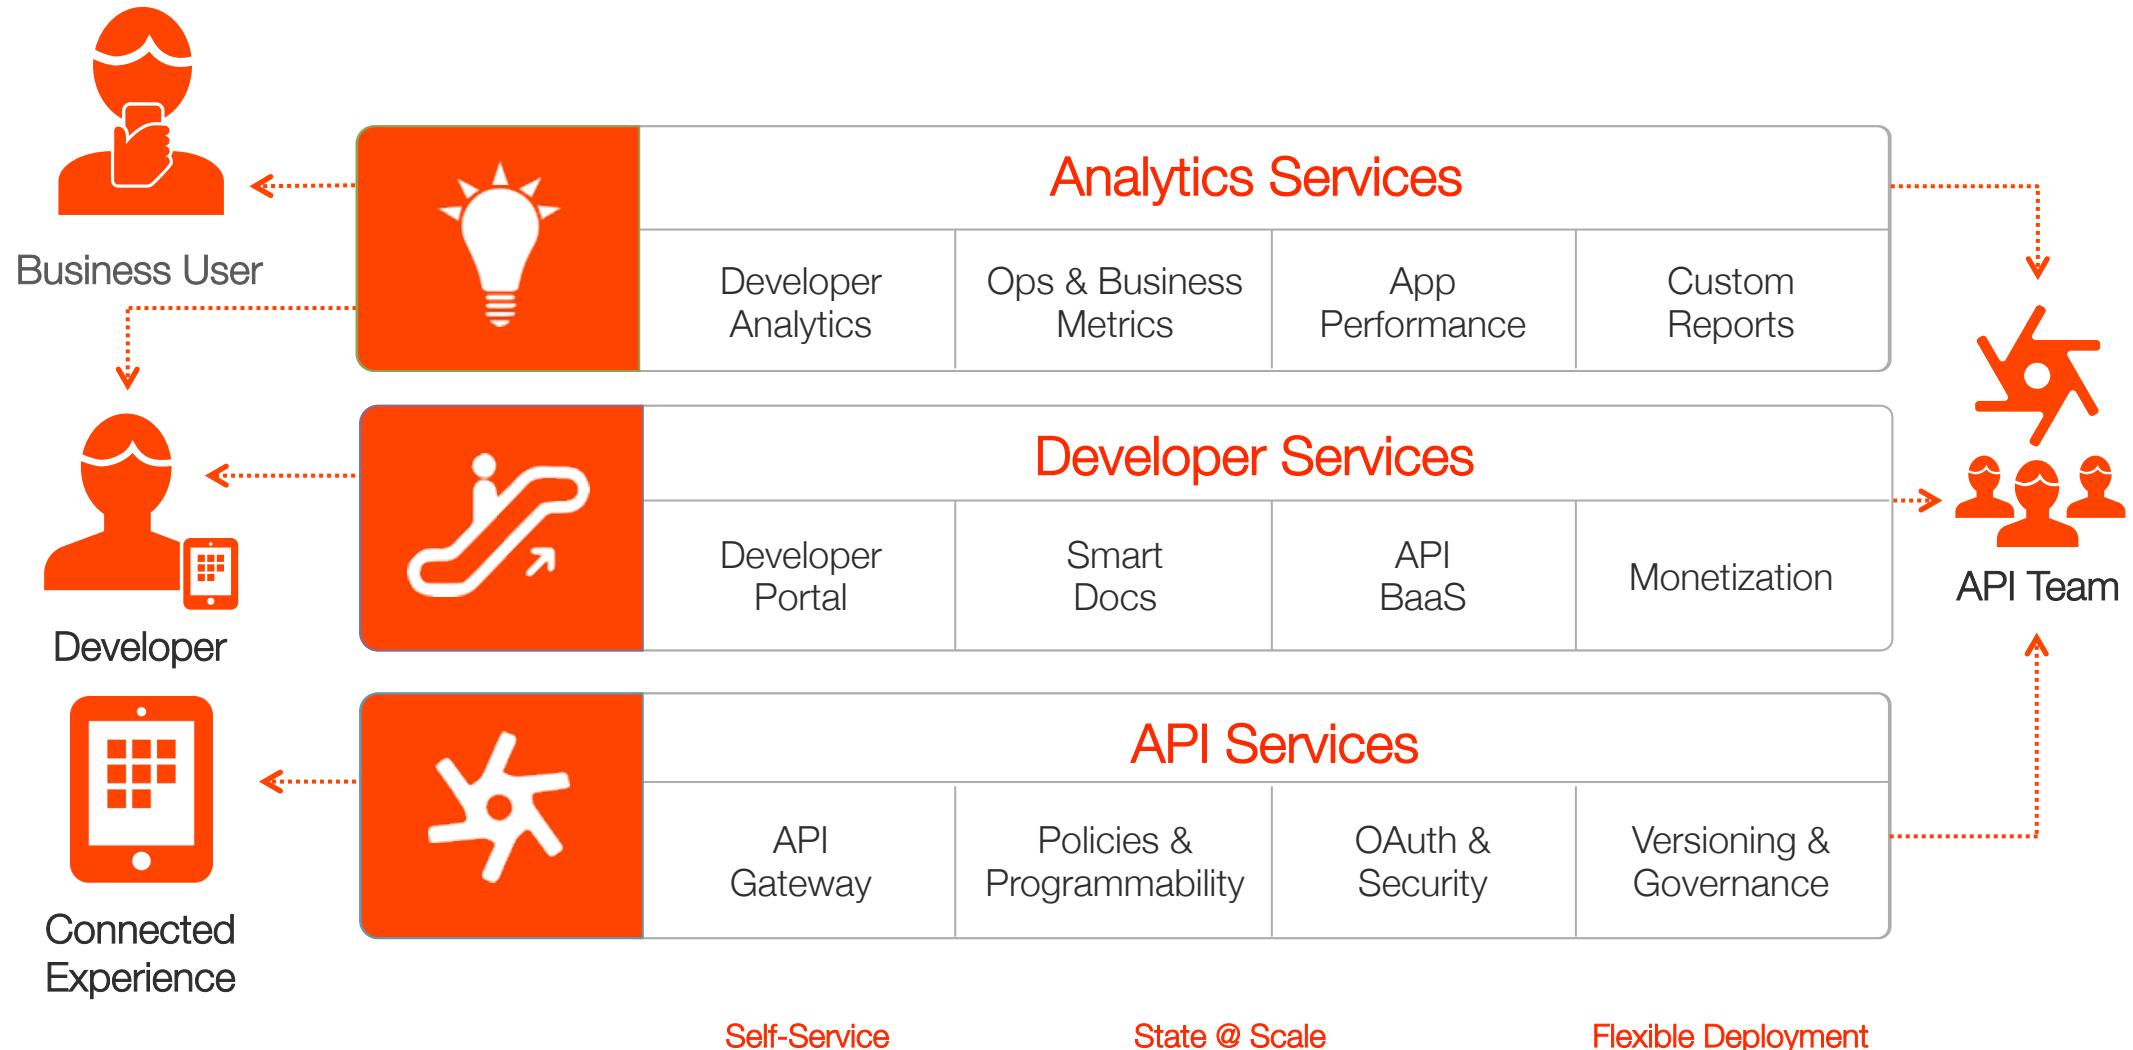
Intelligent API platform enables the digital value chain



Apigee Edge



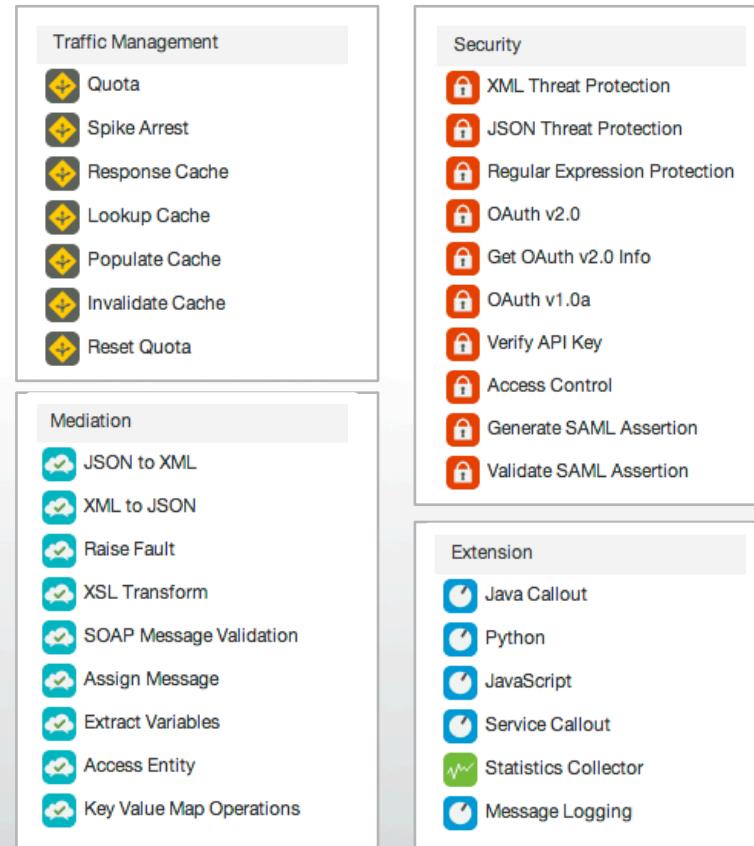
Apigee Edge – API management



Apigee Edge policies – Build APIs faster

Manage interactions with API consumers and optimize performance

Transform, translate and reformat data for easy consumption



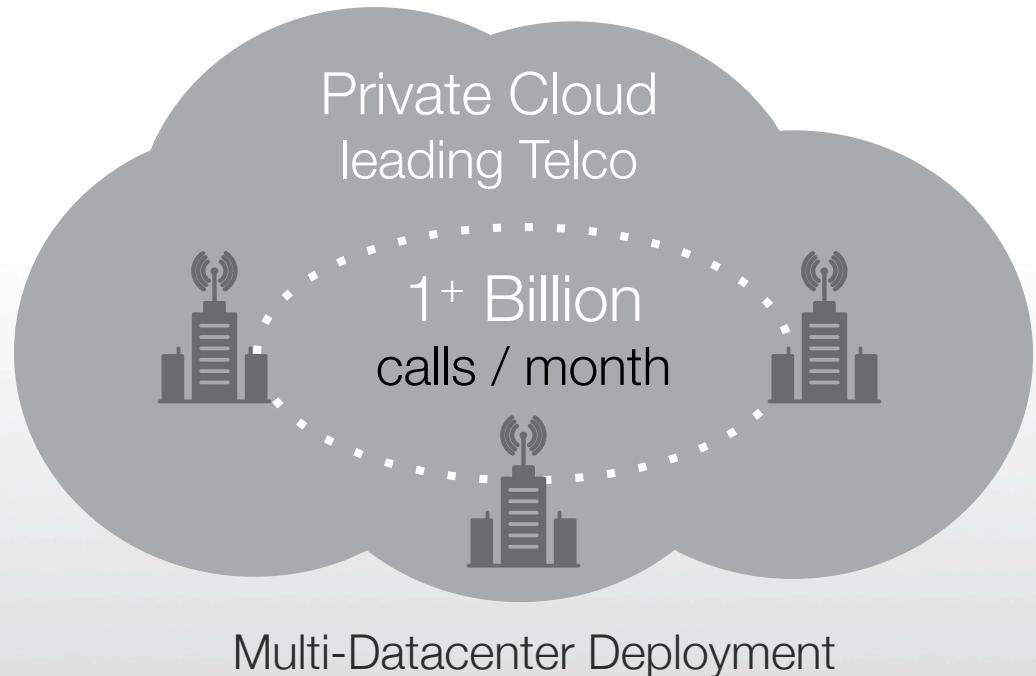
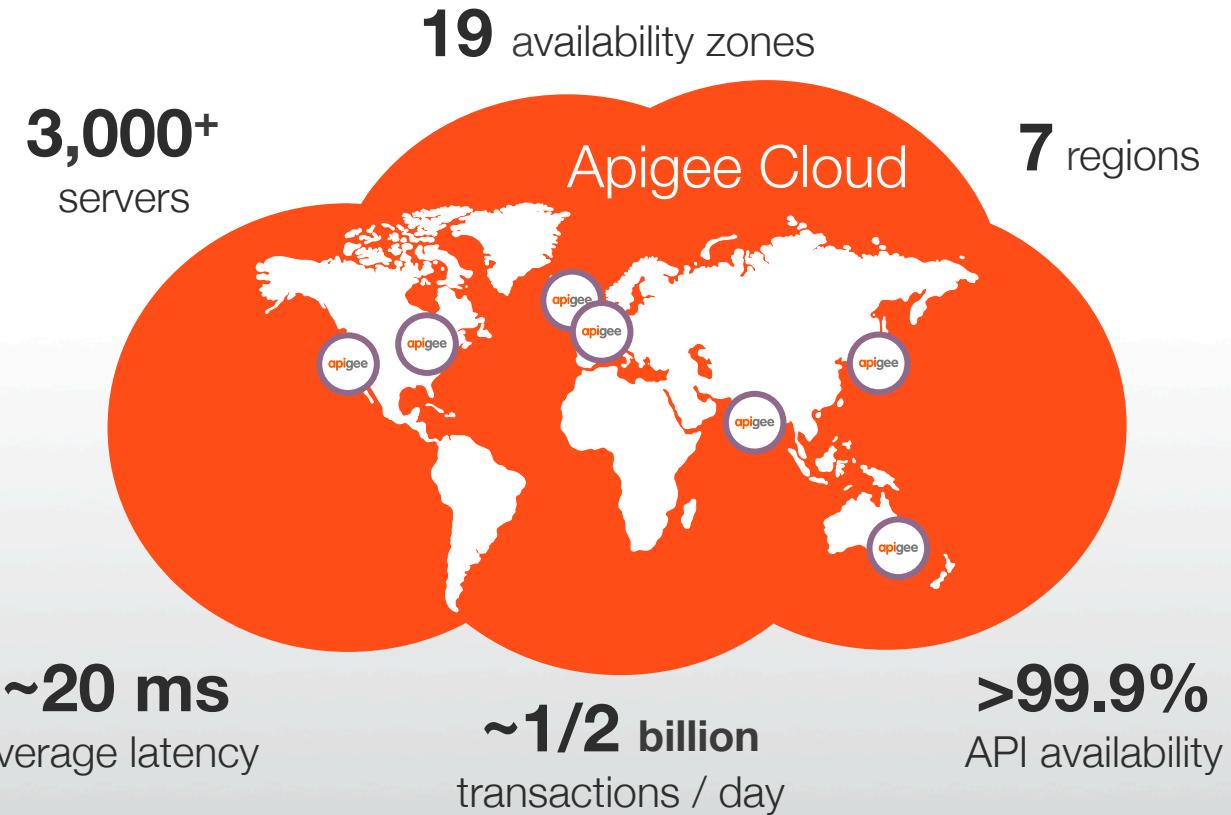
Secure APIs and protect back-end systems from attack

Extend with programming when you need it

Config - Over 30 ready-to-use and configurable policies

Code - Built-in support for Node, JavaScript, Java and Python extensibility

Flexible deployment options



Public Cloud = Private Cloud

Agenda

09:30 – 10:00

Apigee Overview

10:00 – 11:00

Apigee Private Cloud Architecture

BREAK (15 MINS)

11:15 – 12:00

Apigee Private Cloud Architecture (cont)

BREAK – LUNCH (60 MINS)

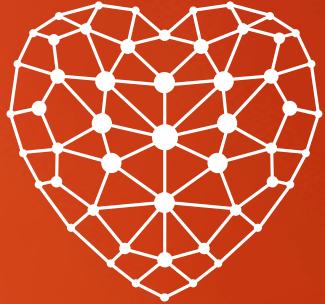
13:00 – 15:45

Apigee Edge Install

BREAK (15 MINS)

16:00 – 17:30

Platform Operations



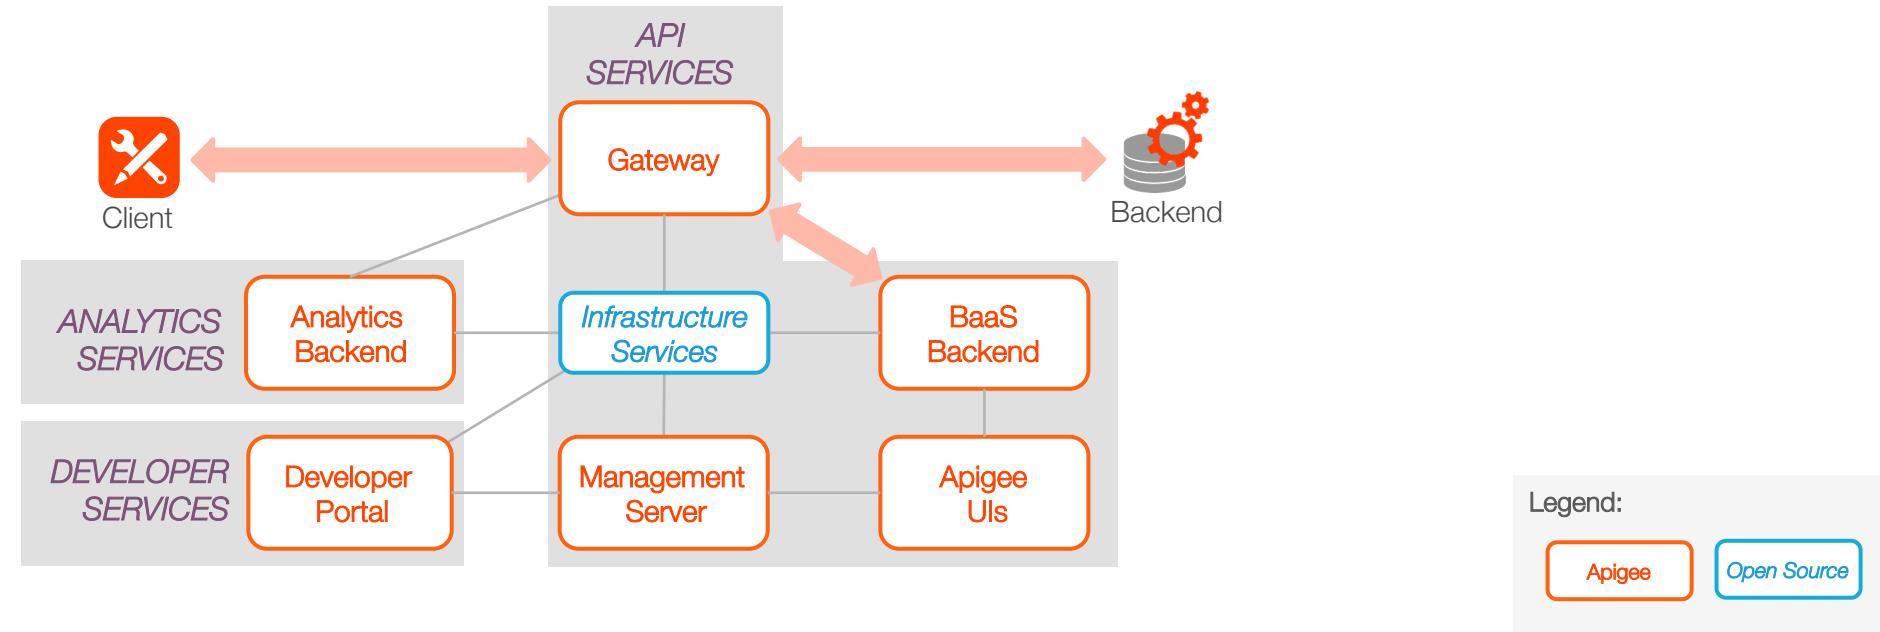
Apigee Private Cloud Architecture

Apigee Edge architecture – High level

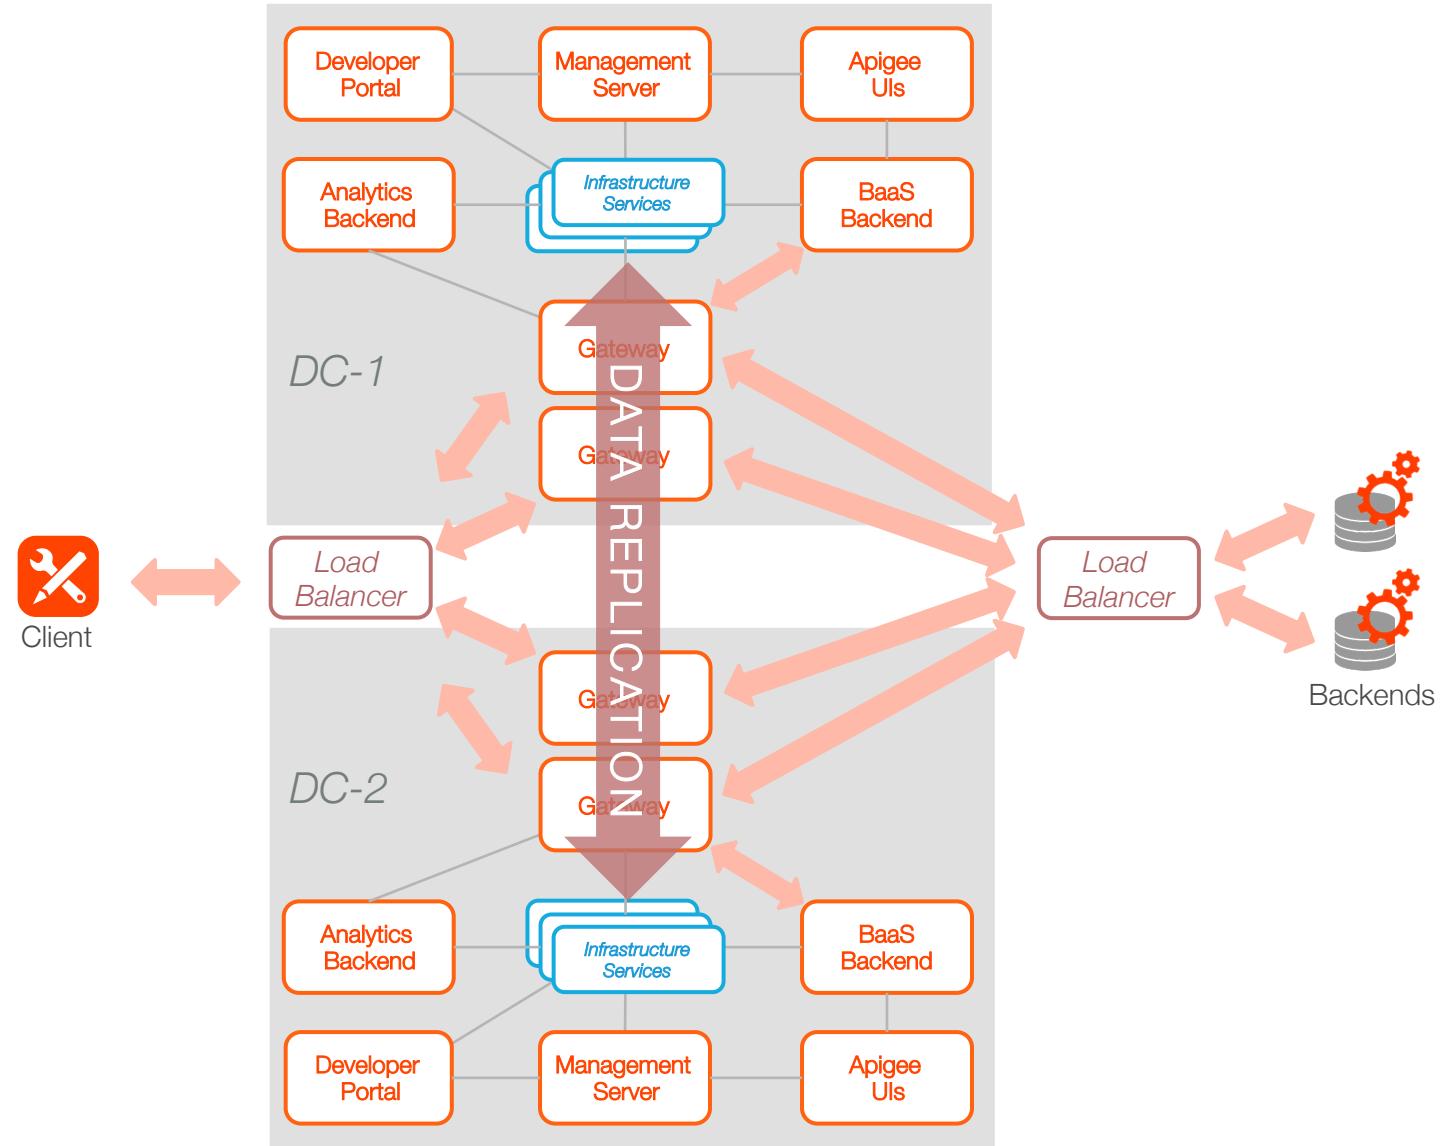
Apigee Edge is comprised of several stateless components that use infrastructure services to persist data

- Gateway: Routing and processing API calls
- Apigee UIs: Enterprise UI, Developer Portal
- Infrastructure Services: Persistence and queuing of run time data
- Management Server: Provider of REST APIs for all configuration tasks

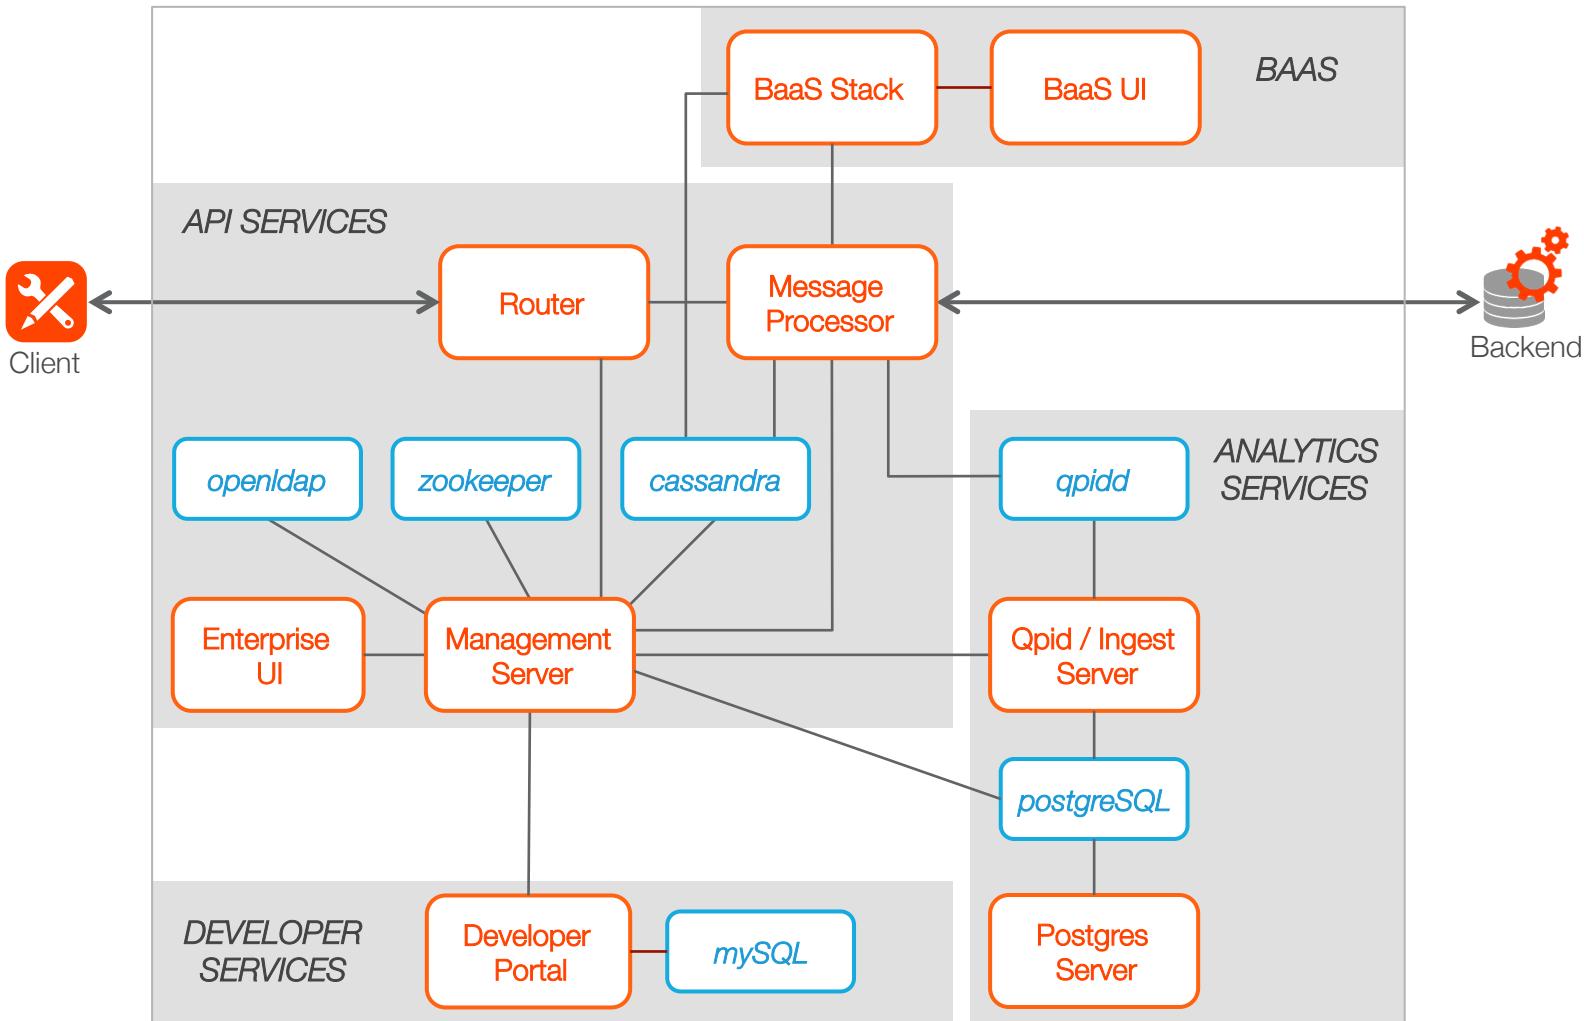
Note: Monetization is part of Developer Services and leverages Gateway, Analytics Services and Management Server



Apigee Edge architecture – High level



Apigee Edge architecture – Component view



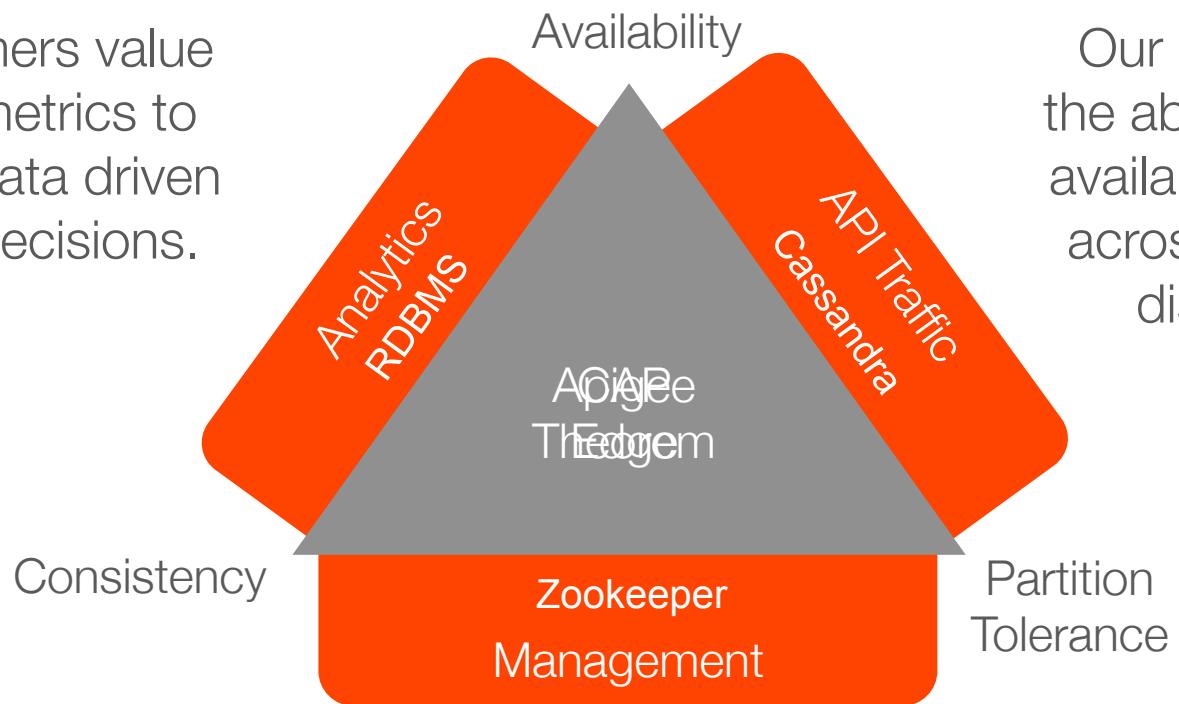
- **Router** handles all incoming API traffic and dispatches it. The Router terminates the HTTP request, handles the SSL traffic, and uses the virtual host name, port, and URI to steer requests to the appropriate node.
- **Message Processor** handles API traffic for a specific organization and environment and which executes all policies.
- Apache **Cassandra** stores application configurations, distributed quota counters, API keys, and OAuth tokens for applications running on the gateway.
- Apache **ZooKeeper** contains configuration data about all the services of the zone and which notifies the different servers of configuration changes.
- **OpenLDAP** contains organization user and roles.
- Management Server offers an API that is used by the Central Services server to communicate with the servers in each on-premises installation.
- **QPID Server** manages queuing system for analytics data.
- **Postgres Server** manages analytics database.

Legend:



Apigee Edge architecture – Technology Stack

Our customers value accurate metrics to help drive data driven business decisions.



Our customers value the ability to have highly available API expanded across geographically dispersed sites.

Our customers value the ability to centralized management of distributed components.

In theoretical computer science, the **CAP theorem**, also known as Brewer's theorem, states that it is impossible for a distributed computer system to simultaneously provide all three of the following guarantees:

- **Consistency** - all nodes see the same data at the same time.
- **Availability** - a guarantee that every request receives a response about whether it succeeded or failed.
- **Partition tolerance** - the system continues to operate despite arbitrary message loss or failure of part of the system.

Text Credits: Wikipedia

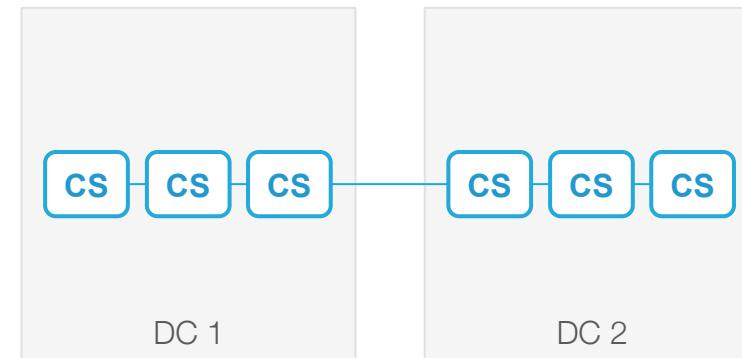
©2015 Apigee. All Rights Reserved.

Apache Cassandra

“Apache Cassandra is an open source distributed database management system. It is an Apache Software Foundation top-level project designed to handle very large amounts of data spread out across many commodity servers while providing a highly available service with no single point of failure.” -- Wikipedia

Cassandra characteristics:

- All nodes are equal. Not master/slave or primary/secondary.
- An application can read/write data from any node.
- Data replication. Apigee Edge uses replication factor 3.
- Consistency managed by application. Apigee Edge uses one and local quorum.



Cassandra is used by Apigee for a variety of purposes, including:

- Storage of developer, application and API Product data
- Storage of access and refresh tokens
- Storage of key-value map data
- Audit logs
- Custom analytics report models

<http://cassandra.apache.org/>

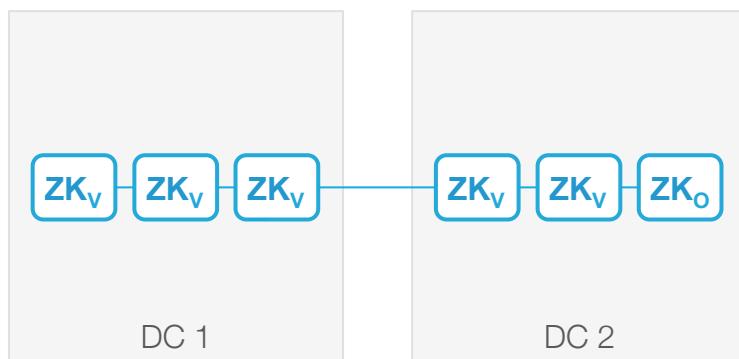
Legend:	R Router	MS Management Server	DP Developer Portal	MY MySQL	OL Openldap	SV Server/Virtual Machine
	MP Message Processor	PS Postgres Server	BA BaaS Server	ZK Zookeeper	PG PostgreSQL	POD POD
	UI Enterprise UI	QIS Qpid/Ingest Server		CS Cassandra	AQ Apache Qpid	

Apache Zookeeper

“Apache ZooKeeper is a software project of the Apache Software Foundation, providing an open source distributed configuration service, synchronization service, and naming registry for large distributed systems.”

– Wikipedia

Zookeeper is used by Apigee as a distributed configuration registry, tracking component location, configuration and status data. With some exceptions, it is NOT required to process API requests.



- **Leader:** The node that controls coordination of writes across distributed Zookeeper nodes
- **Voters:** Nodes that can vote on change proposal made by the Leader
- **Observers:** Do not vote on change proposals and must forward all writes to the Leader

<http://zookeeper.apache.org/>

Legend:

R Router

MP Message Processor

UI Enterprise UI

MS Management Server

PS Postgres Server

QIS Qpid/Ingest Server

DP Developer Portal

BA BaaS Server

CS Cassandra

MY MySQL

ZK Zookeeper

OL Openldap

CS Cassandra

OL Openldap

PG PostgreSQL

QD Apache Qpid

■ Server/Virtual Machine

□ POD

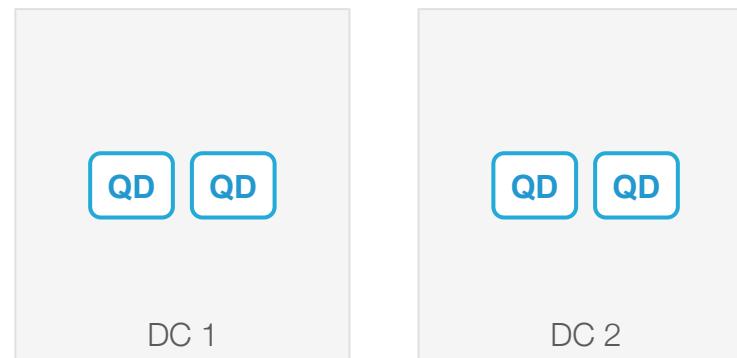
©2015 Apigee. All Rights Reserved.

Apache Qpid

“Apache Qpid, an open-source (Apache 2.0 licensed) messaging system, implements the Advanced Message Queuing Protocol. It provides transaction management, queuing, distribution, security, management, clustering, federation and heterogeneous multi-platform support.”

– Wikipedia

Qpid is used by Apigee Edge as messaging system for analytics and monetization data.



<http://qpid.apache.org/>

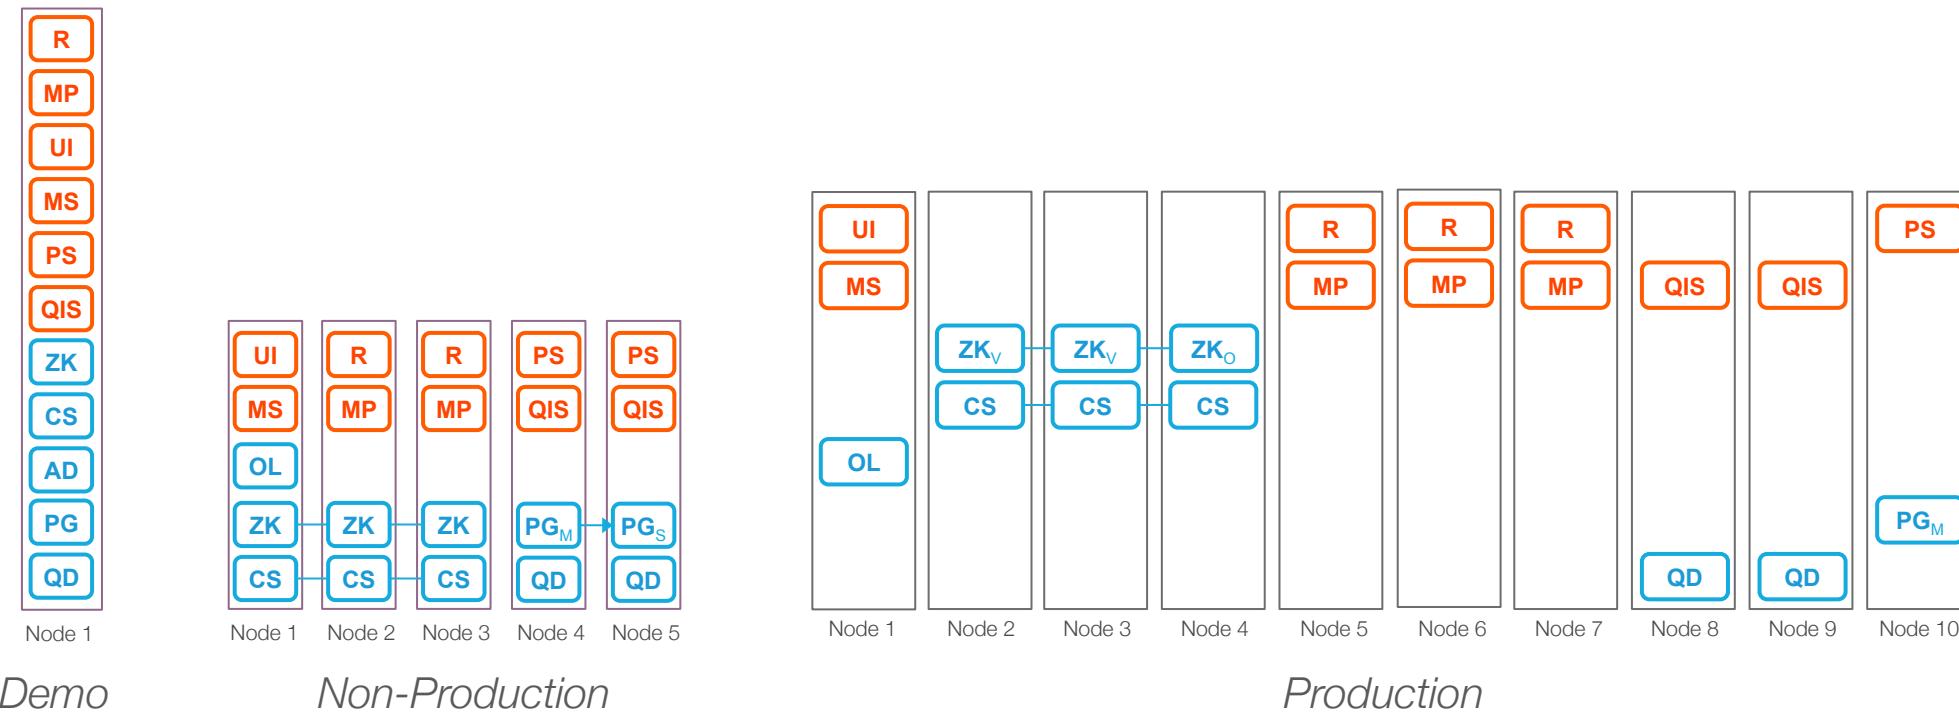
Legend:

R	Router	MS	Management Server	DP	Developer Portal	MY	MySQL	OL	Openldap	PS	Server/Virtual Machine
MP	Message Processor	PS	Postgres Server	BA	BaaS Server	ZK	Zookeeper	PG	PostgreSQL	POD	POD
UI	Enterprise UI	QIS	Qpid/Ingest Server			CS	Cassandra	QD	Apache Qpid		

©2015 Apigee. All Rights Reserved.

Apigee Edge architecture – Deployment options

- Apigee Edge architecture offers great flexibility when it comes to deployment options.
- Edge can be deployed from a single VM to a multi-datacenter active/active configuration.
- Edge has been designed from the ground up to be a true cloud scale solution, capable of running on both virtualized (including AWS) and physical hardware.



Legend:

R Router

MP Message Processor

UI Enterprise UI

MS Management Server

PS Postgres Server

QIS Qpid/Ingest Server

DP Developer Portal

BA BaaS Server

CS Cassandra

MY MySQL

ZK Zookeeper

CS Cassandra

OL Openldap

PG PostgreSQL

QD Apache Qpid

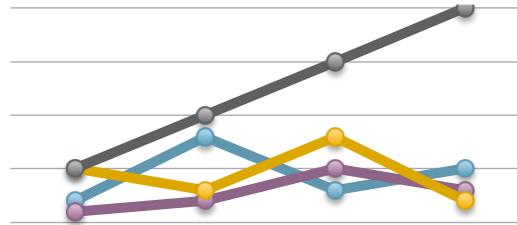
POD Server/Virtual Machine

Topology selection – Key considerations

1. Know your business



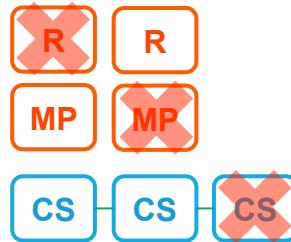
2. Understand traffic patterns



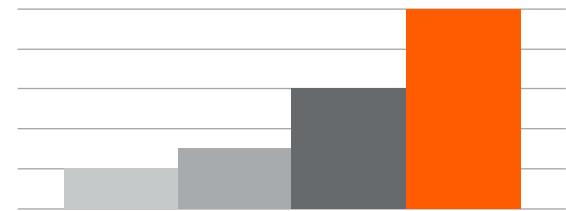
4. Know your API proxies

Traffic Management	Security
Quota	XML Threat Protection
Spike Arrest	JSON Threat Protection
Response Cache	Regular Expression Protection
Lookup Cache	OAuth v2.0
Populate Cache	Get OAuth v2.0 Info
Invalidate Cache	OAuth v1.0a
Reset Quota	Verify API Key
Mediation	Access Control
JSON to XML	Generate SAML Assertion
XML to JSON	Validate SAML Assertion
Raise Fault	
XSL Transform	
SOAP Message Validation	
Assign Message	
Extract Variables	
Access Entity	
Key Value Map Operations	Statistics Collector
Extension	Message Logging
Java Callout	
Python	
JavaScript	
Service Callout	
Statistics Collector	
Message Logging	

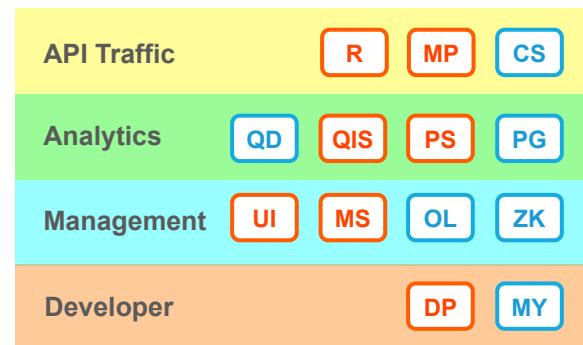
5. Remember everything fails



3. Plan for growth (2X, 5X, 10X?)



6. API Traffic, Analytics and Management, Developer Service components can and should scale independently.



Legend:

R Router

MS Management Server

DP Developer Portal

MY MySQL

OL Openldap

Server/Virtual Machine

MP Message Processor

PS Postgres Server

BA BaaS Server

ZK Zookeeper

PG PostgreSQL

POD

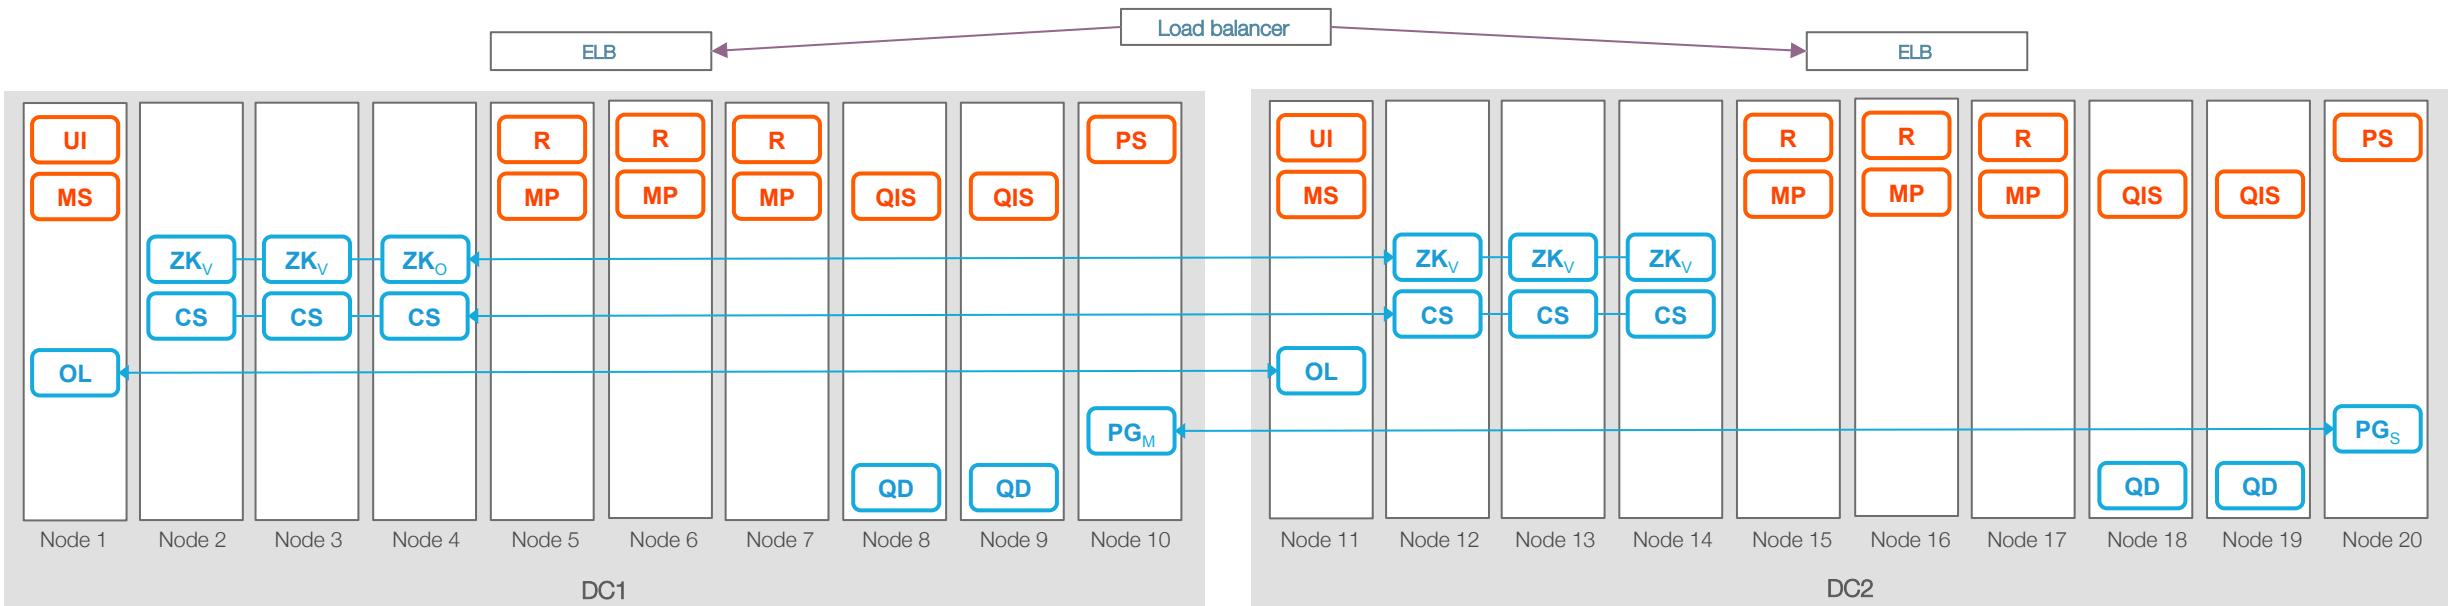
UI Enterprise UI

QIS Qpid/Ingest Server

CS Cassandra

Apache Qpid

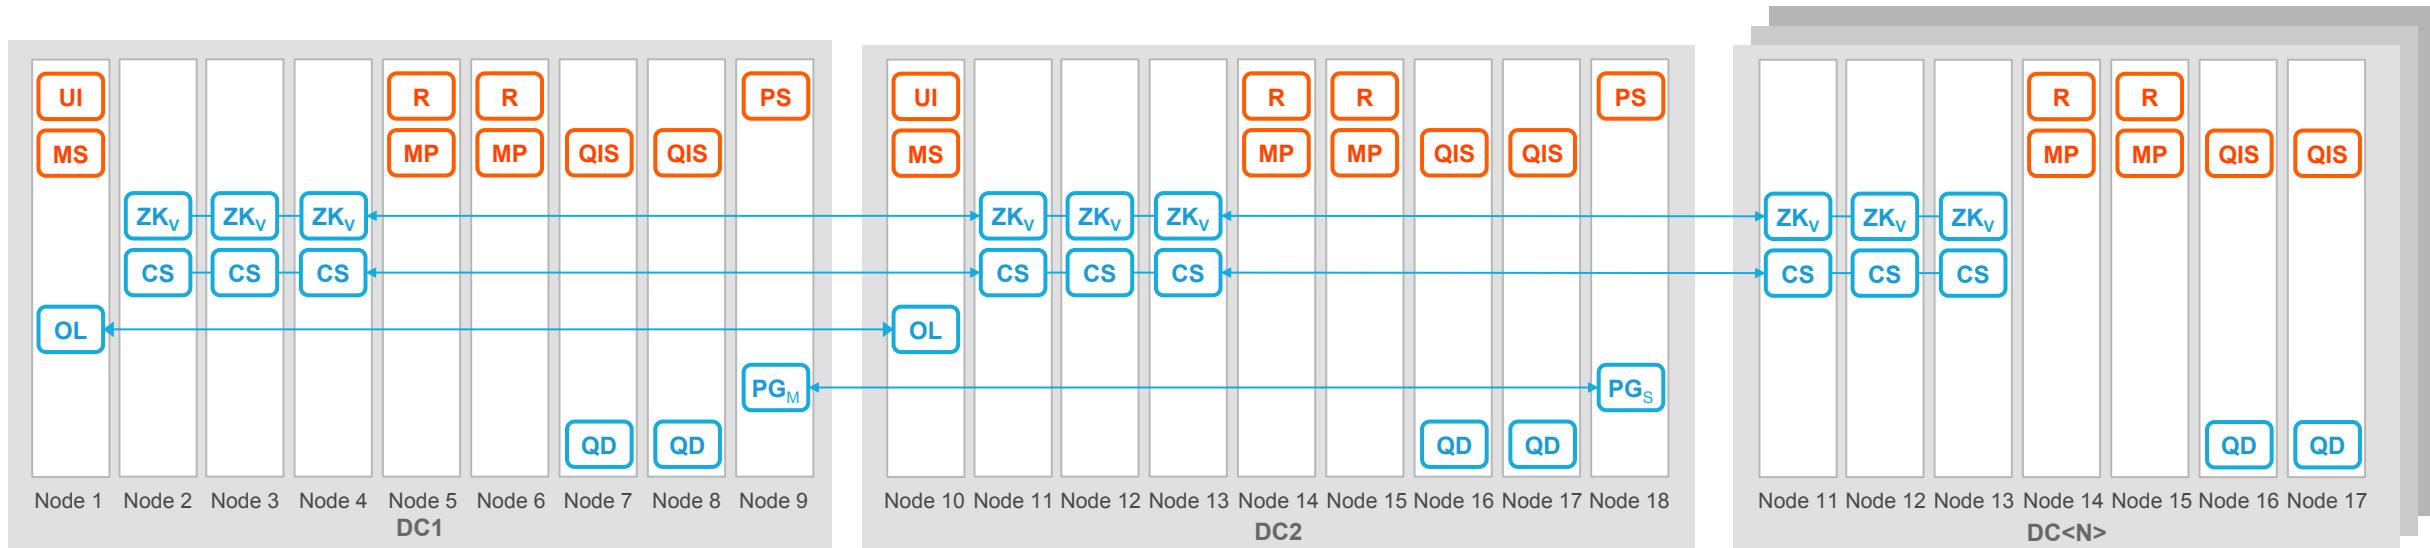
Apigee Edge architecture – Architecture characteristics



- Horizontally scalable. Components can be installed and configured to run from a single node (all in one setup) all the way to a multi DC centers, active/active, globally distributed setup.
- Ability to stack components.
- Distributed data replicated using eventual consistency.
- Asynchronous analytics data capture and processing.
- Centralized configuration for distributed components.
- Multitenant by design.
- Management via APIs and UI console.

Legend:	R Router	MS Management Server	DP Developer Portal	MY MySQL	OL Openldap	Server/Virtual Machine
	MP Message Processor	PS Postgres Server	BA BaaS Server	ZK Zookeeper	PG PostgreSQL	POD
	UI Enterprise UI	QIS Qpid/Ingest Server		CS Cassandra	QD Apache Qpid	

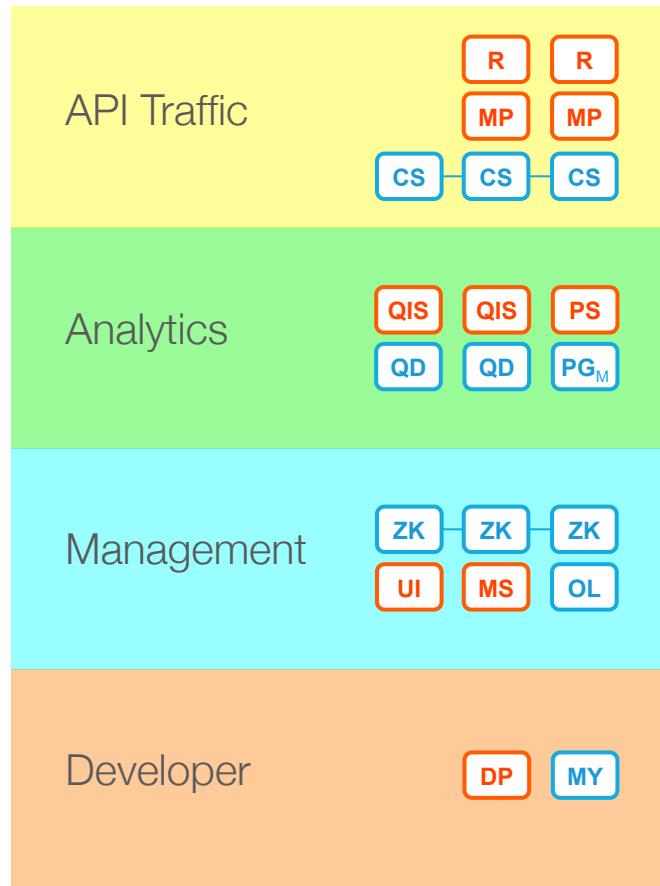
Apigee Edge architecture – Horizontal scalability



- Planet footprint is driven by customer requirements, transaction volumes, availability and reliability among others drive component stacking and number of nodes.
- R, MP and CS are critical components to handle live API traffic.
- Management and analytics components are not required on every DC.
- For high availability, ZK and QD should be available in all DCs.
- Analytics data replication via Master/Slave or Publish/Subscribe.

Legend:	R Router	MS Management Server	DP Developer Portal	MY MySQL	OL Openldap	Server/Virtual Machine
	MP Message Processor	PS Postgres Server	BA BaaS Server	ZK Zookeeper	PG PostgreSQL	POD
	UI Enterprise UI	QIS Qpid/Ingest Server		CS Cassandra	QD Apache Qpid	

Apigee Edge architecture – Scaling by capability



- Given the responsibility and capabilities offered by each component, scalability requirements and how they are implemented may vary.
- In most scenarios, scaling to accommodate higher number of TPS or API calls may impact only components serving live API traffic.
- Analytics and management components may grow in number mostly driven by demanding high availability requirements for those capabilities provided by analytics and management components.

Legend:

Router

Message Processor

Enterprise UI

Management Server

Postgres Server

Qpid/Ingest Server

Developer Portal

BaaS Server

Cassandra

MySQL

Zookeeper

Apache Qpid

Openldap

PostgreSQL

POD

Server/Virtual Machine



15 mins break

Agenda

09:30 – 10:00

Apigee Overview

10:00 – 11:00

Apigee Private Cloud Architecture

BREAK (15 MINS)

11:15 – 12:00

Apigee Private Cloud Architecture (cont)

BREAK – LUNCH (60 MINS)

13:00 – 15:45

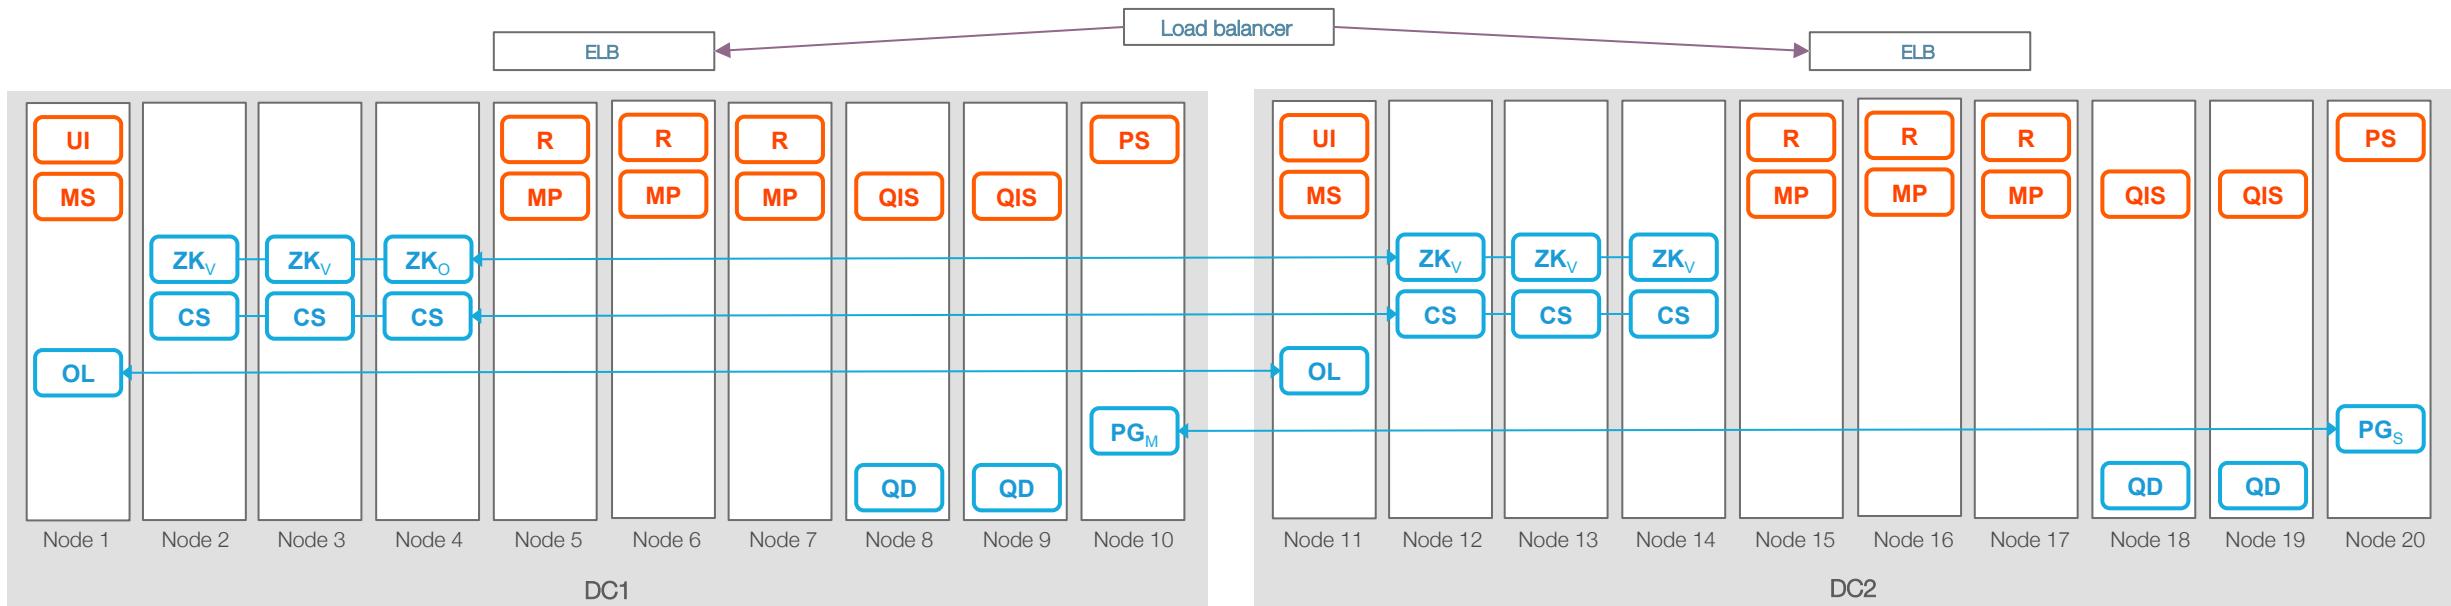
Apigee Edge Install

BREAK (15 MINS)

16:00 – 17:30

Platform Operations

Apigee Edge architecture – Multitenant by design



Legend:

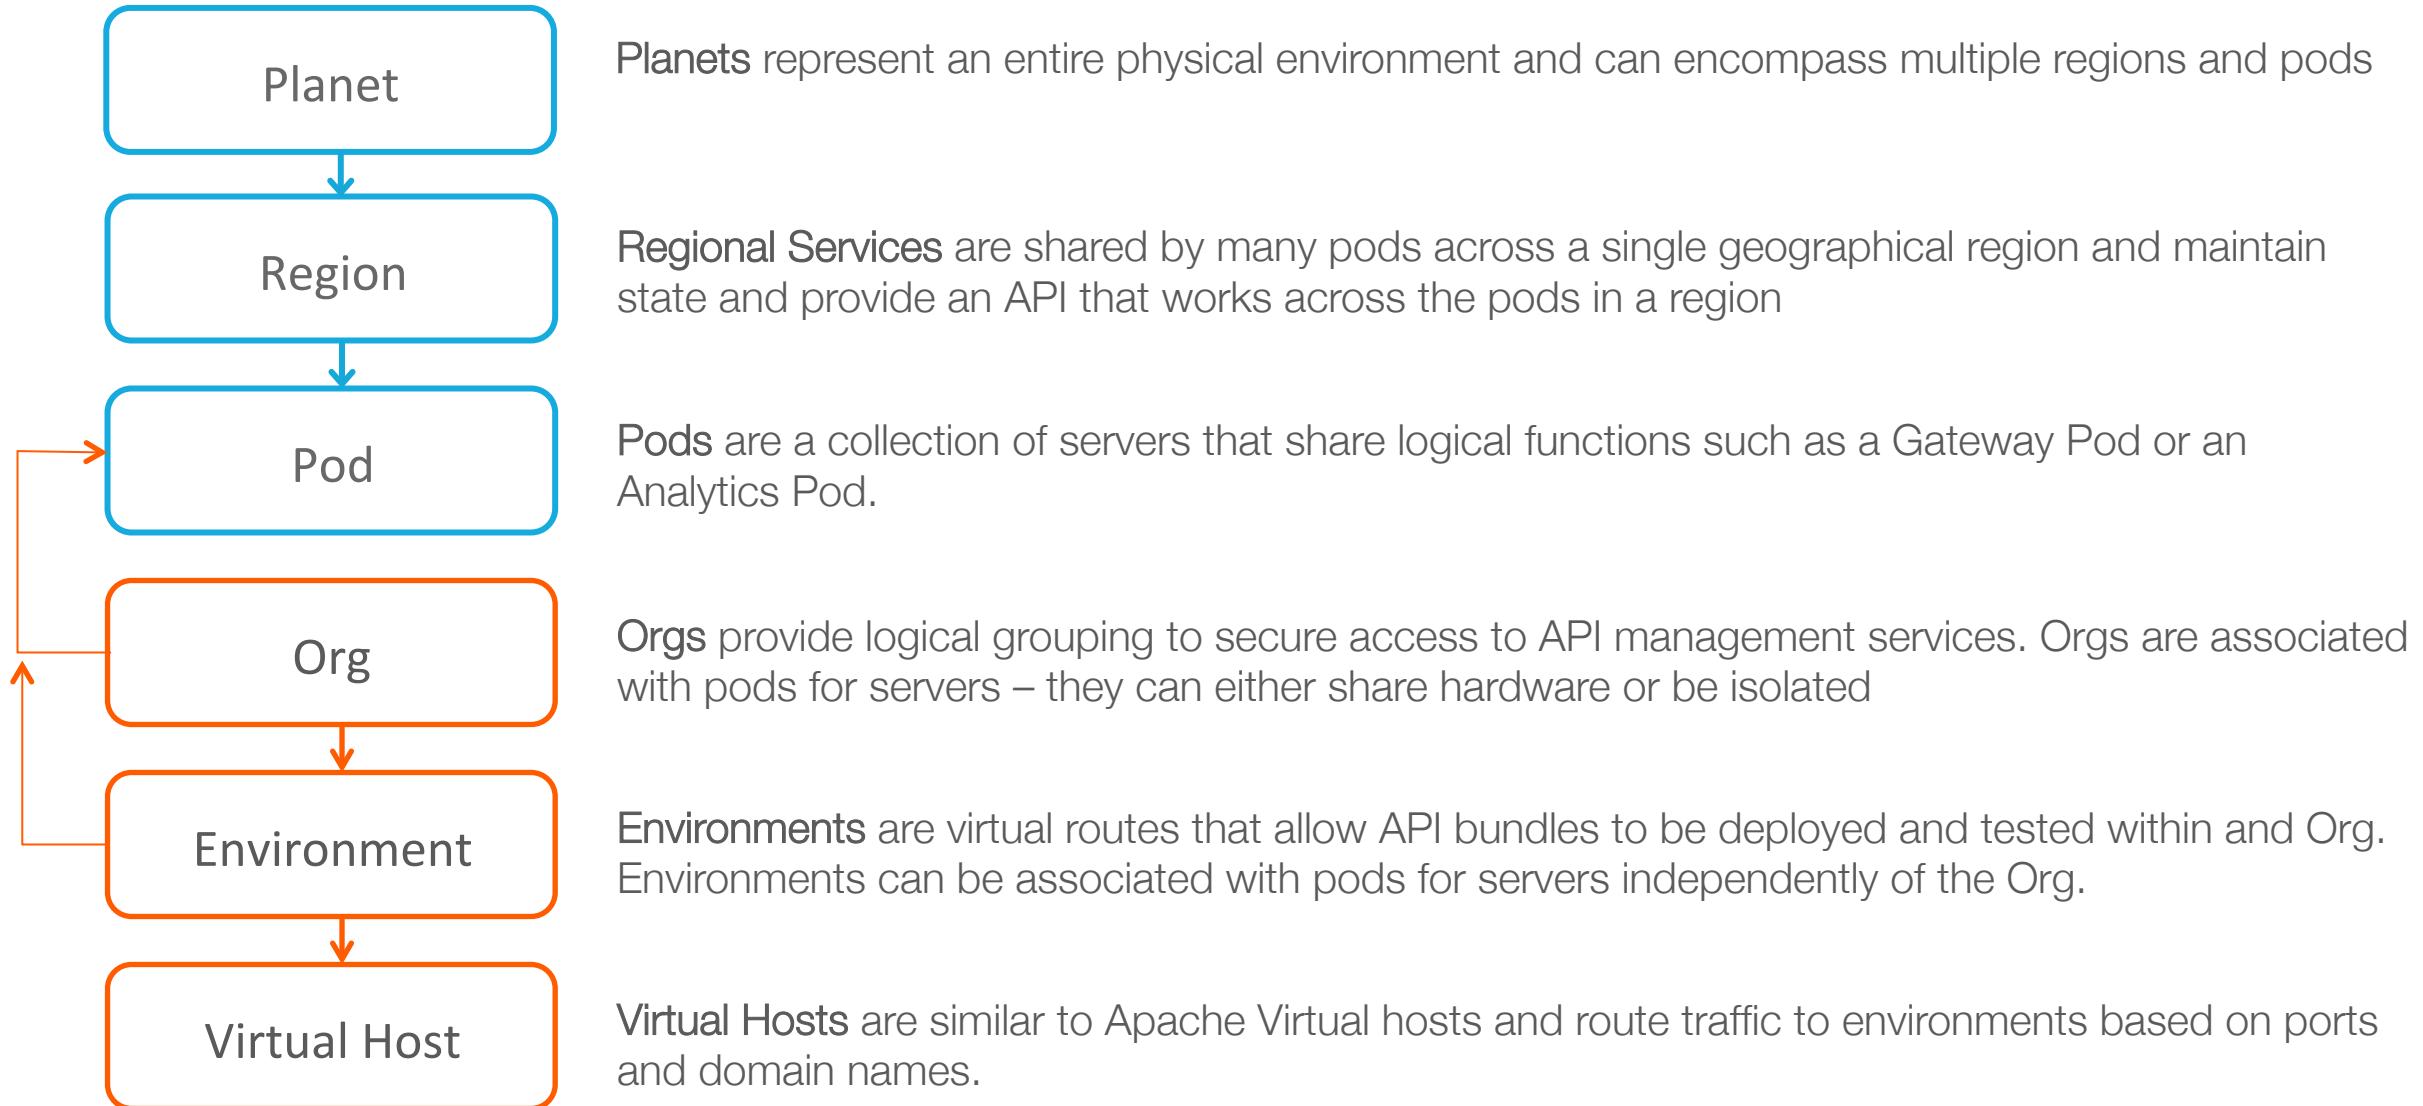
R	Router	MS	Management Server	DP	Developer Portal	MY	MySQL	OL	Openldap	Server/Virtual Machine
MP	Message Processor	PS	Postgres Server	BA	BaaS Server	ZK	Zookeeper	PG	PostgreSQL	POD
UI	Enterprise UI	QIS	Qpid/Ingest Server			CS	Cassandra	QD	Apache Qpid	

Multitenancy

“Multitenancy is a reference to the mode of operation of software where multiple independent instances of one or multiple applications operate in a shared environment. The instances (tenants) are logically isolated, but physically integrated. The degree of logical isolation must be complete, but the degree of physical integration will vary. The more physical integration, the harder it is to preserve the logical isolation. The tenants (application instances) can be representations of organizations that obtained access to the multitenant application (this is the scenario of an ISV offering services of an application to multiple customer organizations). The tenants may also be multiple applications competing for shared underlying resources (this is the scenario of a private or public cloud where multiple applications are offered in a common cloud environment).” - Gartner

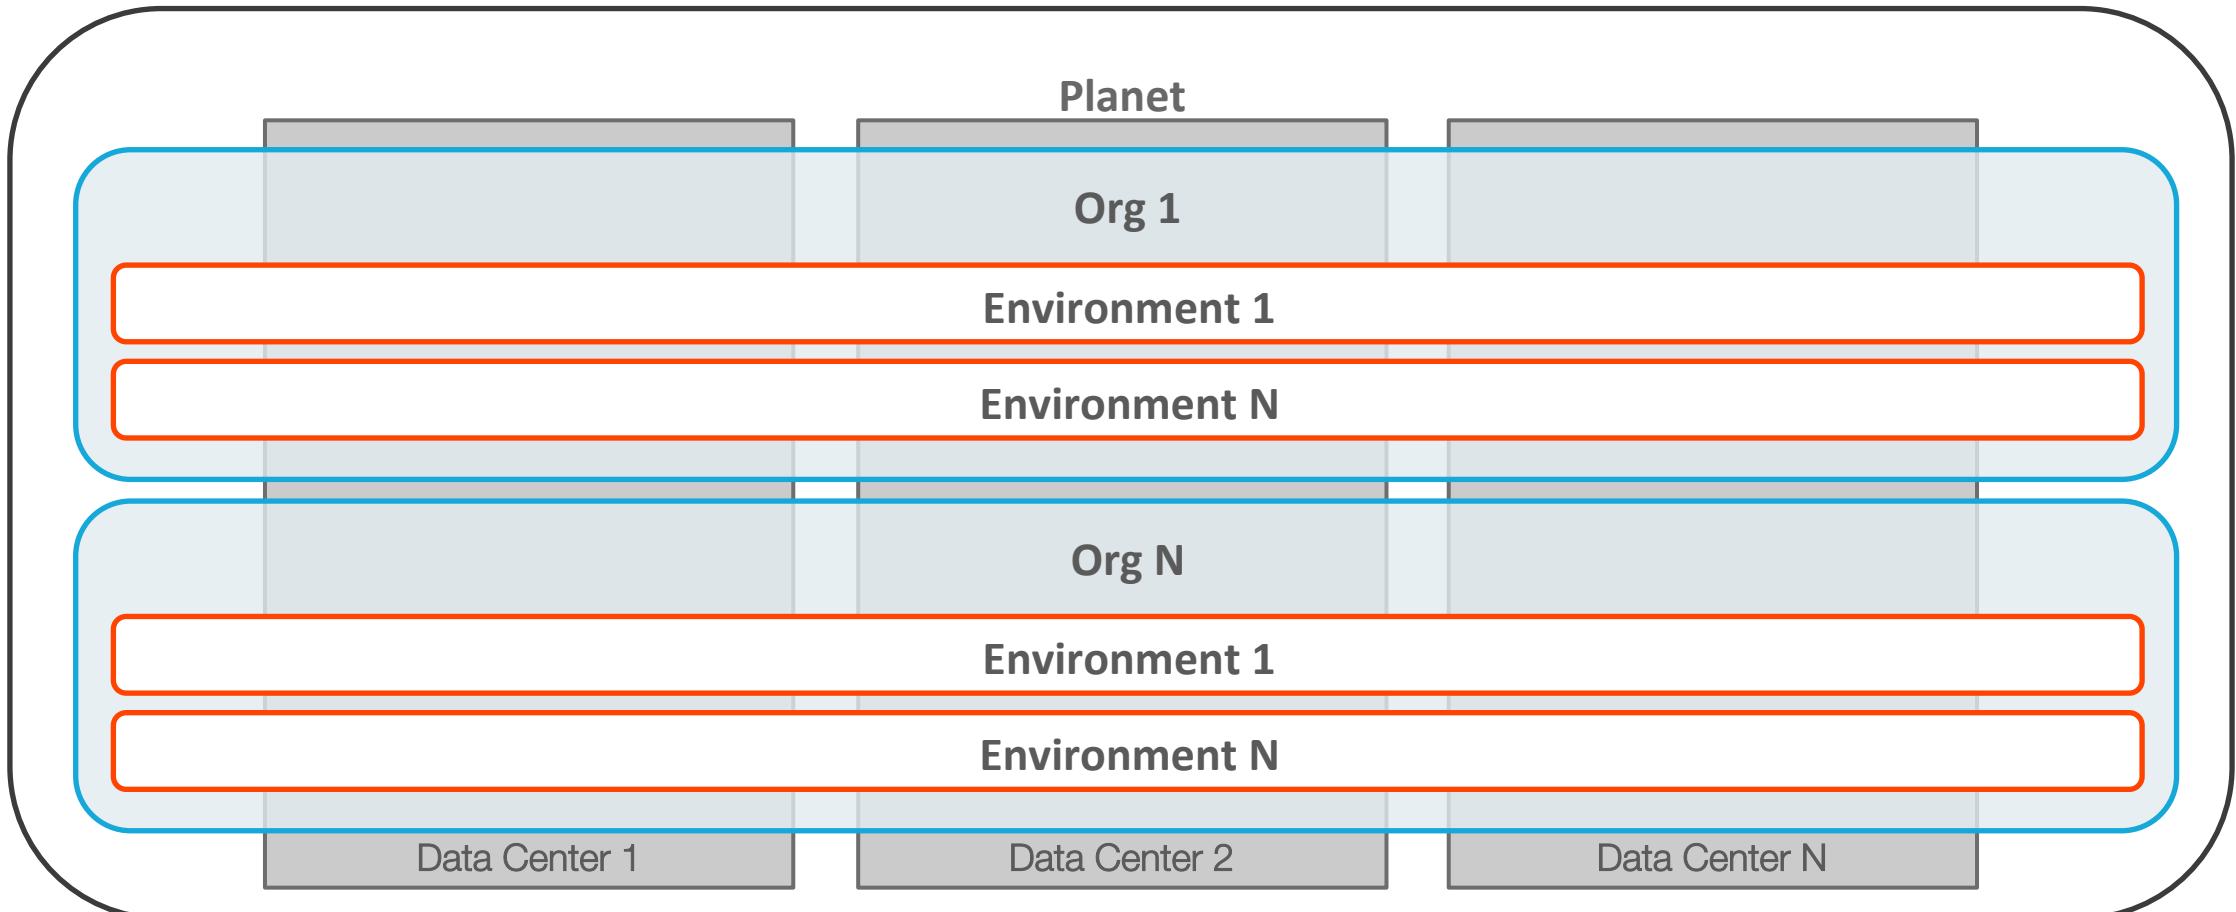
Multitenancy = Software capabilities to support shared infrastructure while providing tenants with data and processing partitioning.

Apigee Edge Architecture – Multitenancy

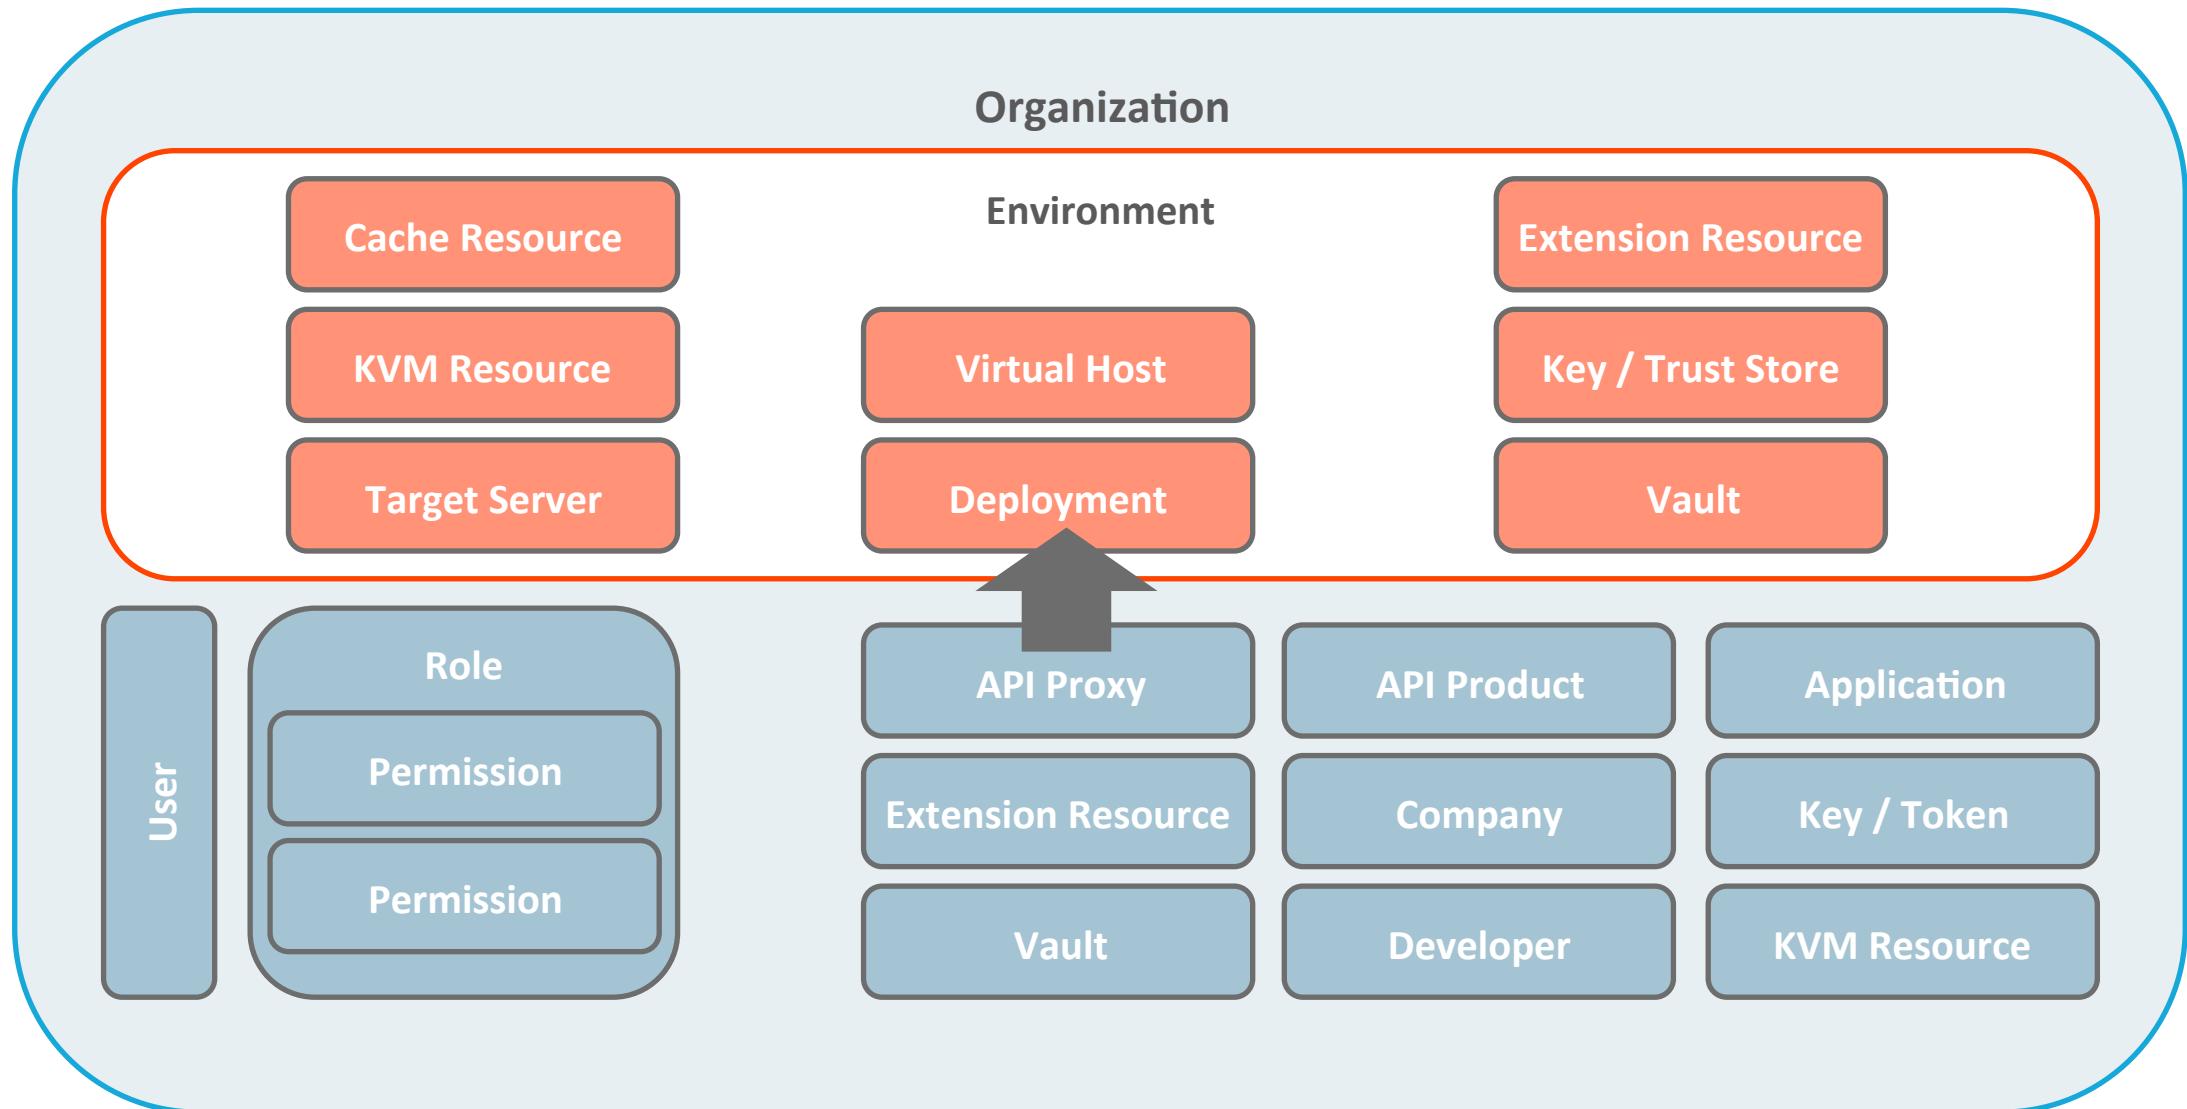


Apigee Edge Architecture – Multitenancy

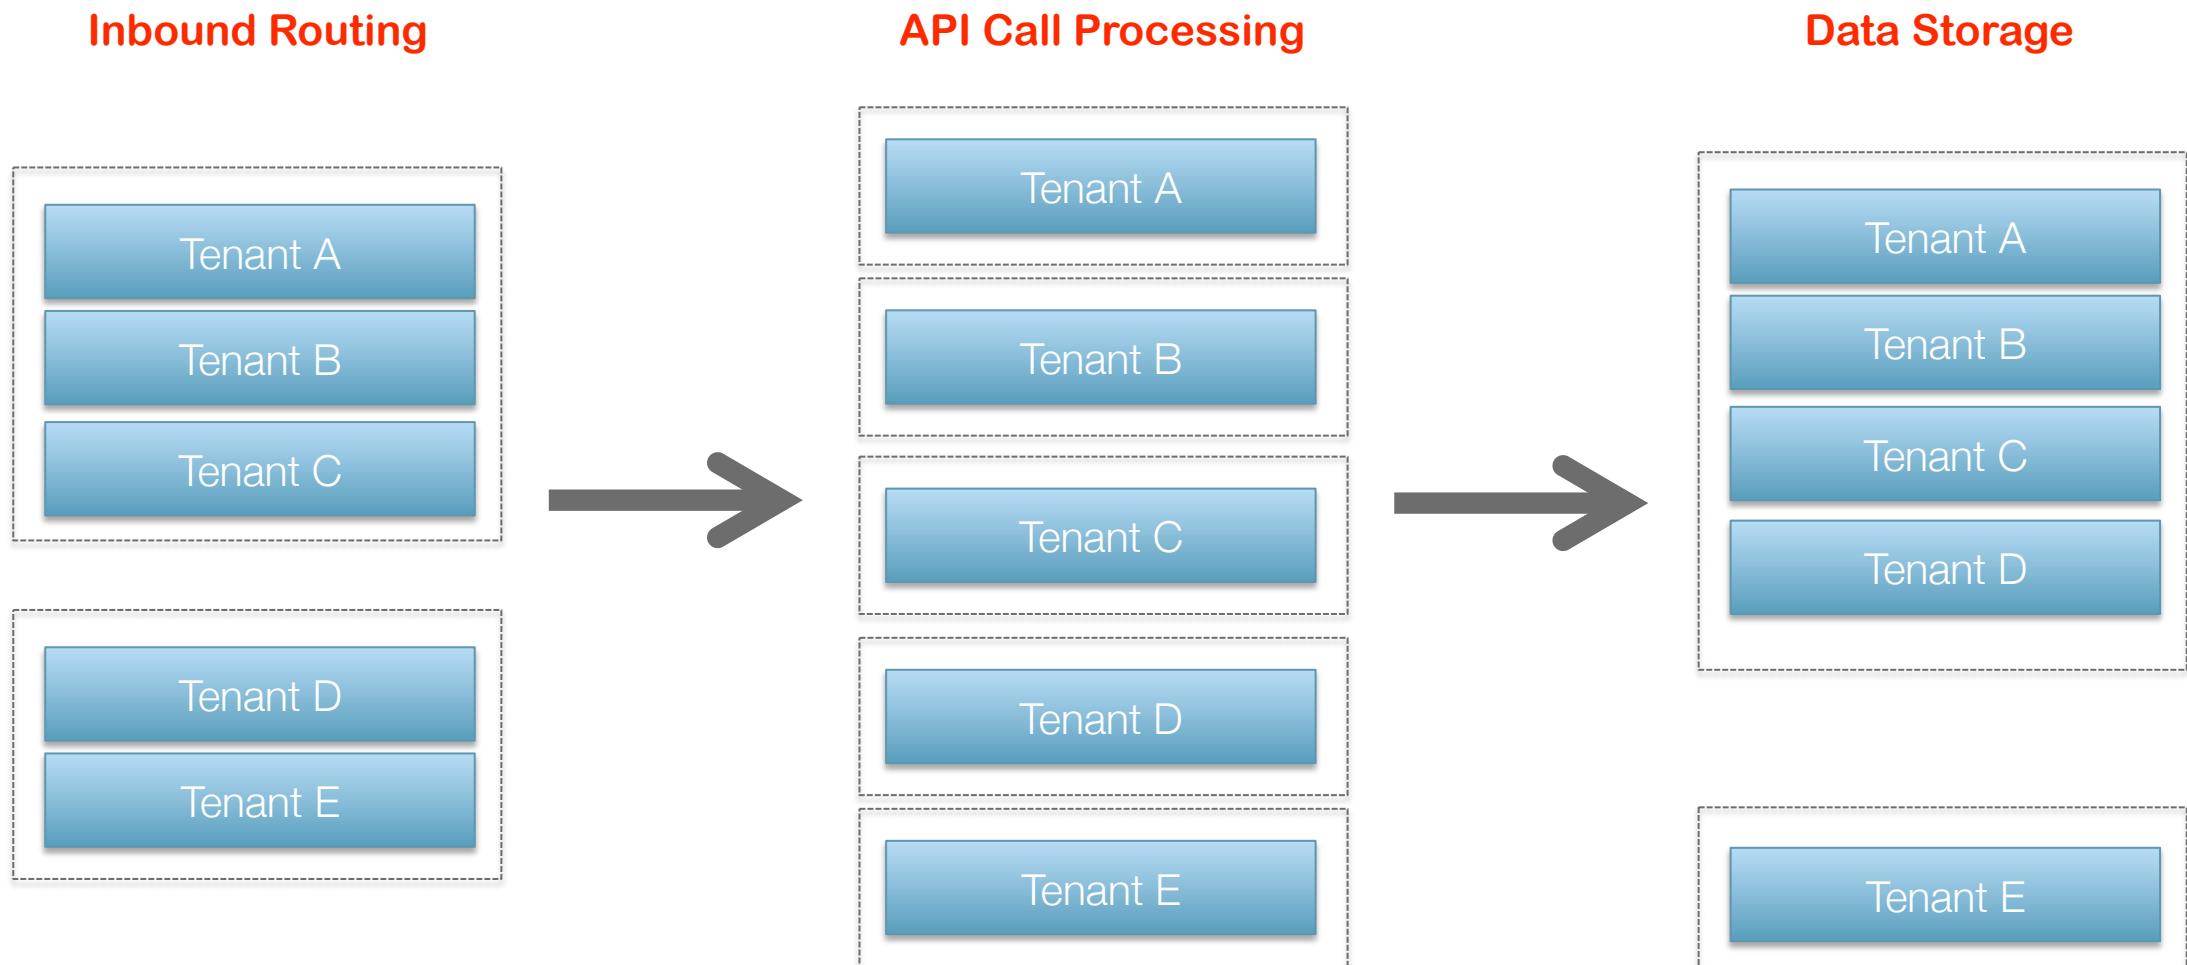
- A Planet can expand multiple DCs. Organization and Environment can expand across the planet.
- Data partitioning by Organization and Environment. Processing partitioning by Org+Env can be configured.



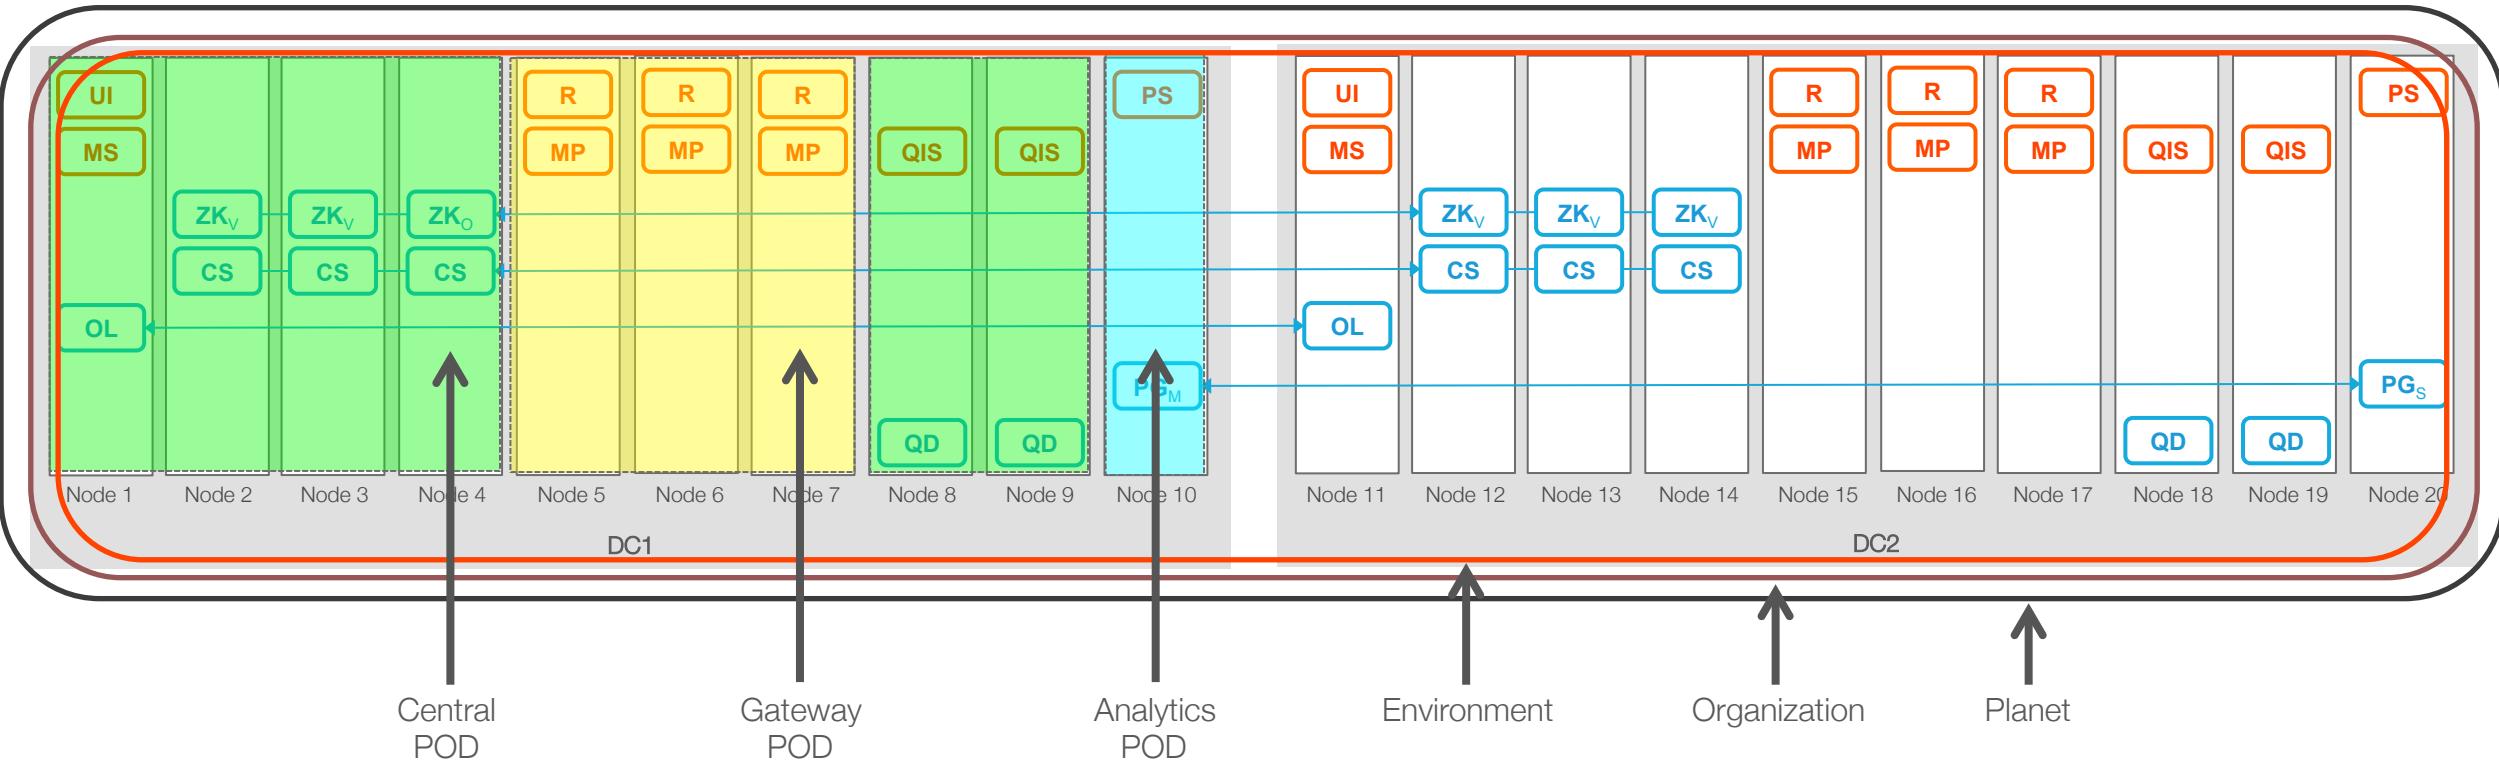
Apigee Edge Architecture – Organizational structure



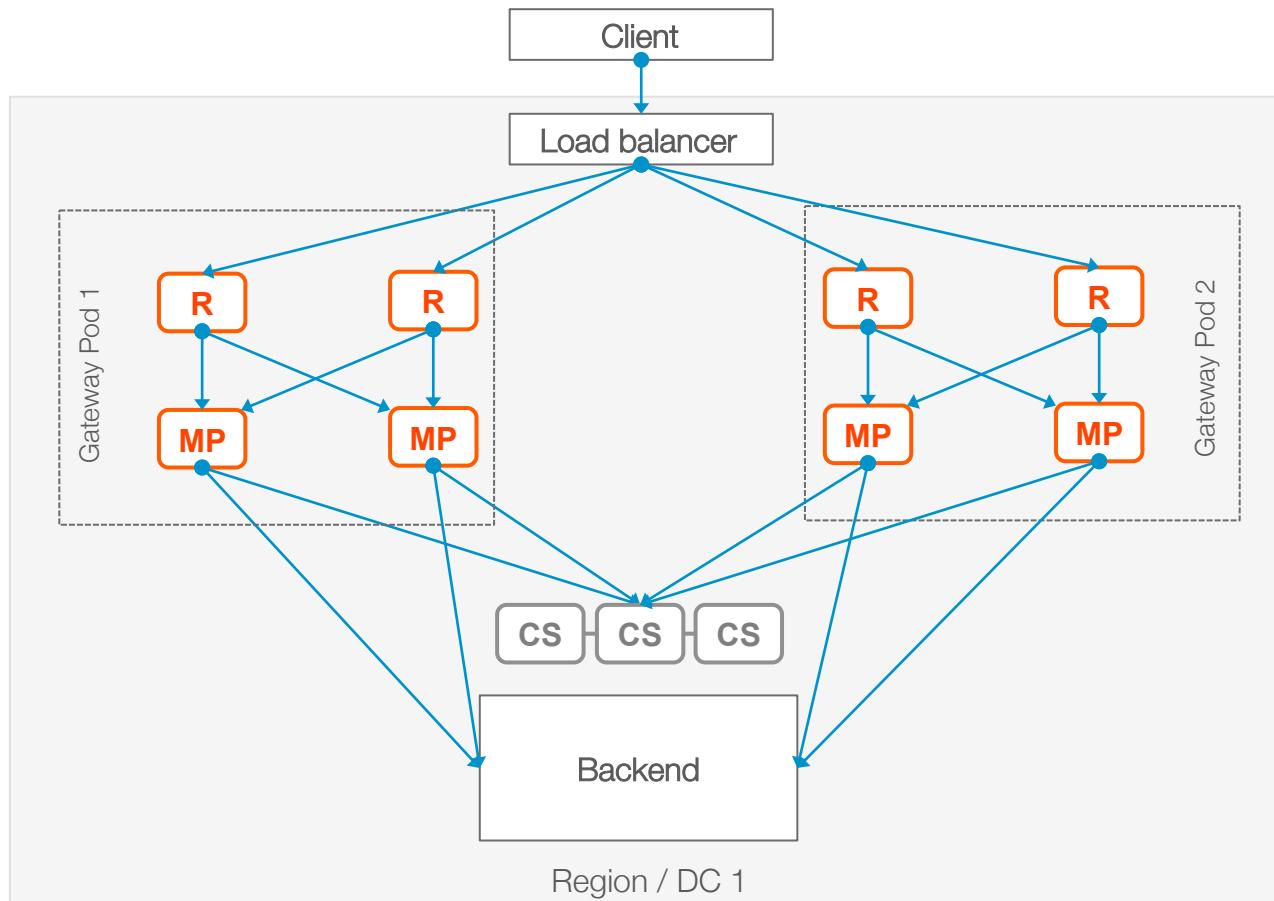
Apigee Edge Architecture – Physical partitioning



Apigee Edge Architecture – Organizational structure



Apigee Edge Architecture – API traffic data flow

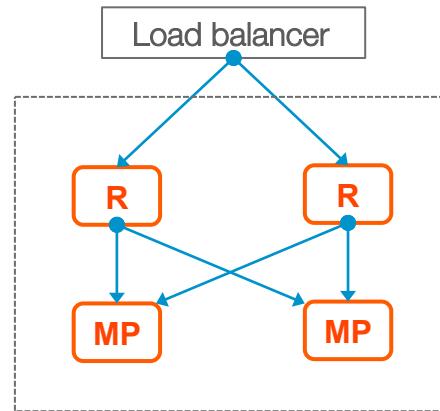


- Routers send requests to Message Processors in their Gateway pod.
- If there are two or more gateway pods in a region, then routers will ignore message processors in the other gateway pods.
- Message Processors respect the region as their scope.
- For two data center the same rules apply as for one datacenter.
- All Apigee components are configured to only use the Cassandra nodes in their region / data center.

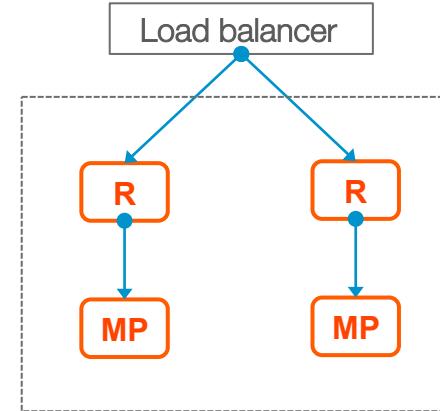
Apigee Edge Architecture – API traffic data flow

R+MP Configuration options

Tenant Aware Routing



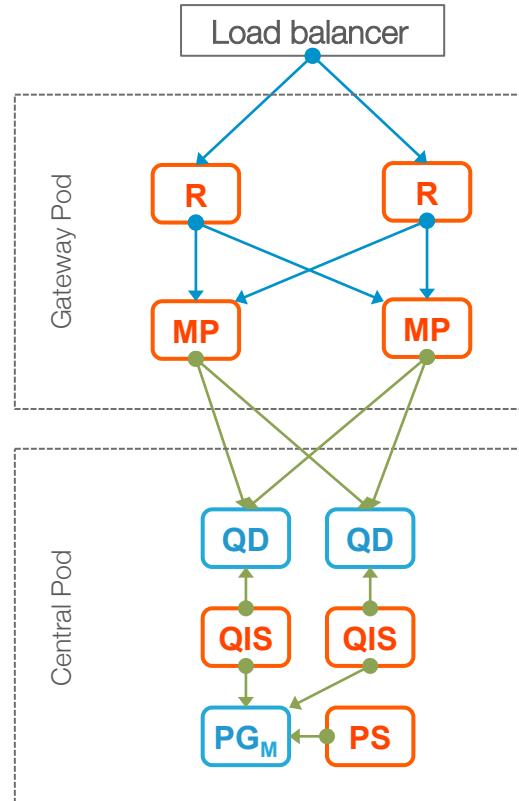
Server Affinity



- Rs direct traffic to appropriate MPs, load balancing between them.
- Default behavior.
- Health check heartbeat allows R to automatically take MP out of/into rotation if unresponsiveness is detected.

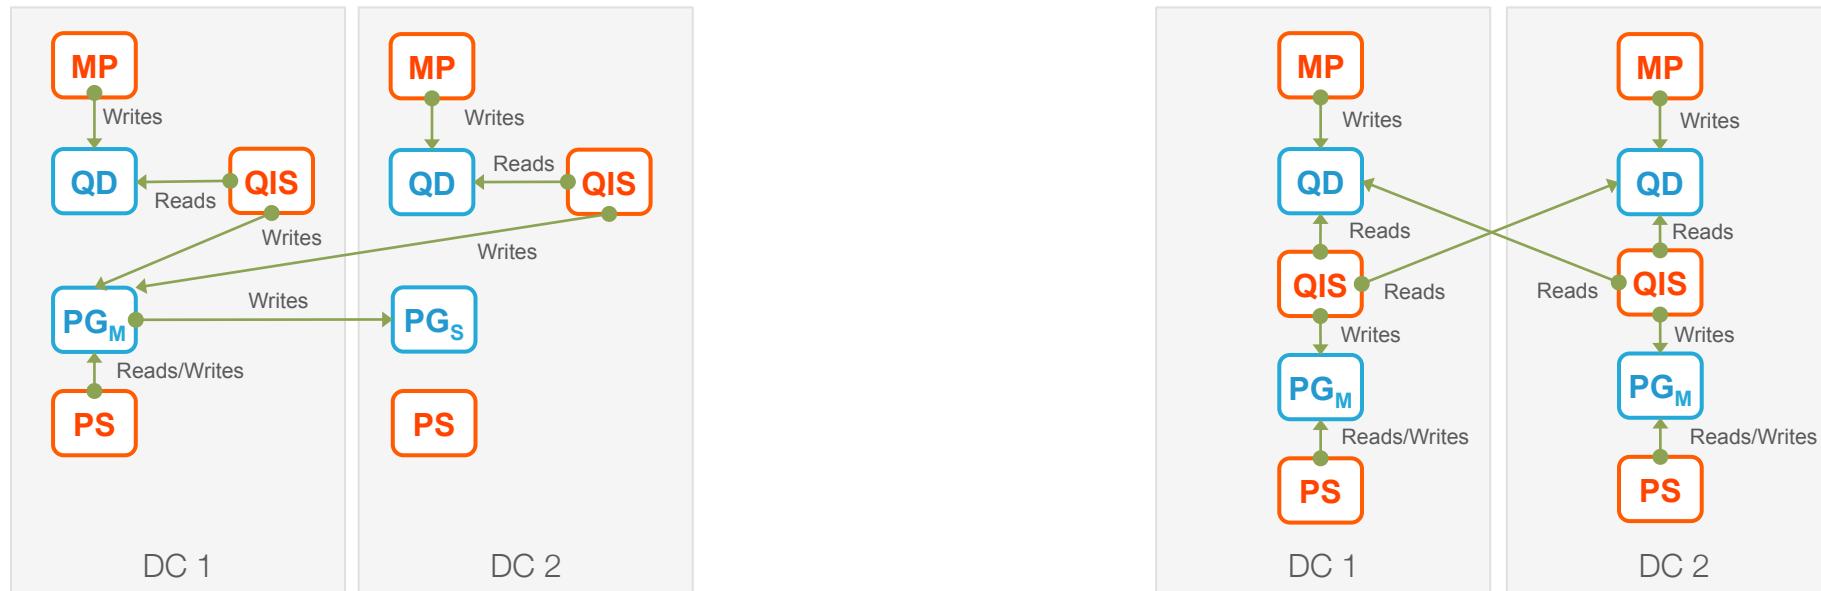
- R can also be configured to connect to a dedicated MP. Using server affinity, all traffic handled by a R is exclusively sent to its corresponding MP.
- This configuration option offers customer the ability to isolate R/MP for dedicated use cases without impacting API traffic flowing across other R/MP within the same pod.

Apigee Edge Architecture – Analytics data flow



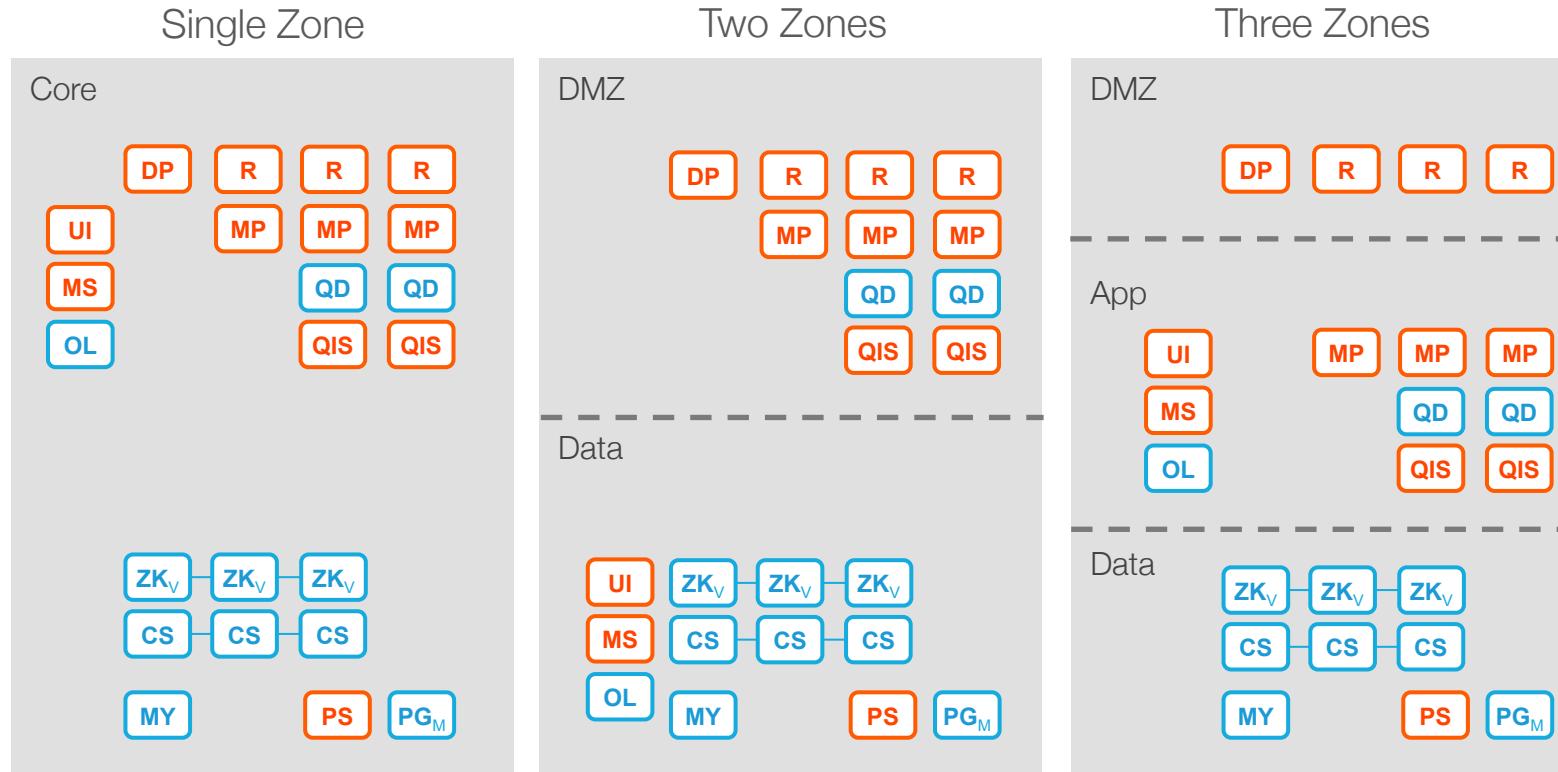
- Ingest services in Qpid Ingest Server will collect Analytics data from all queues and store in PostgreSQL.
- Postgres Server aggregates analytics data asynchronously.
- Message Processors respect the region as their scope and will offload analytics data to all Apache Qpid queues in their region / data center.

Apigee Edge Architecture – Analytics data flow



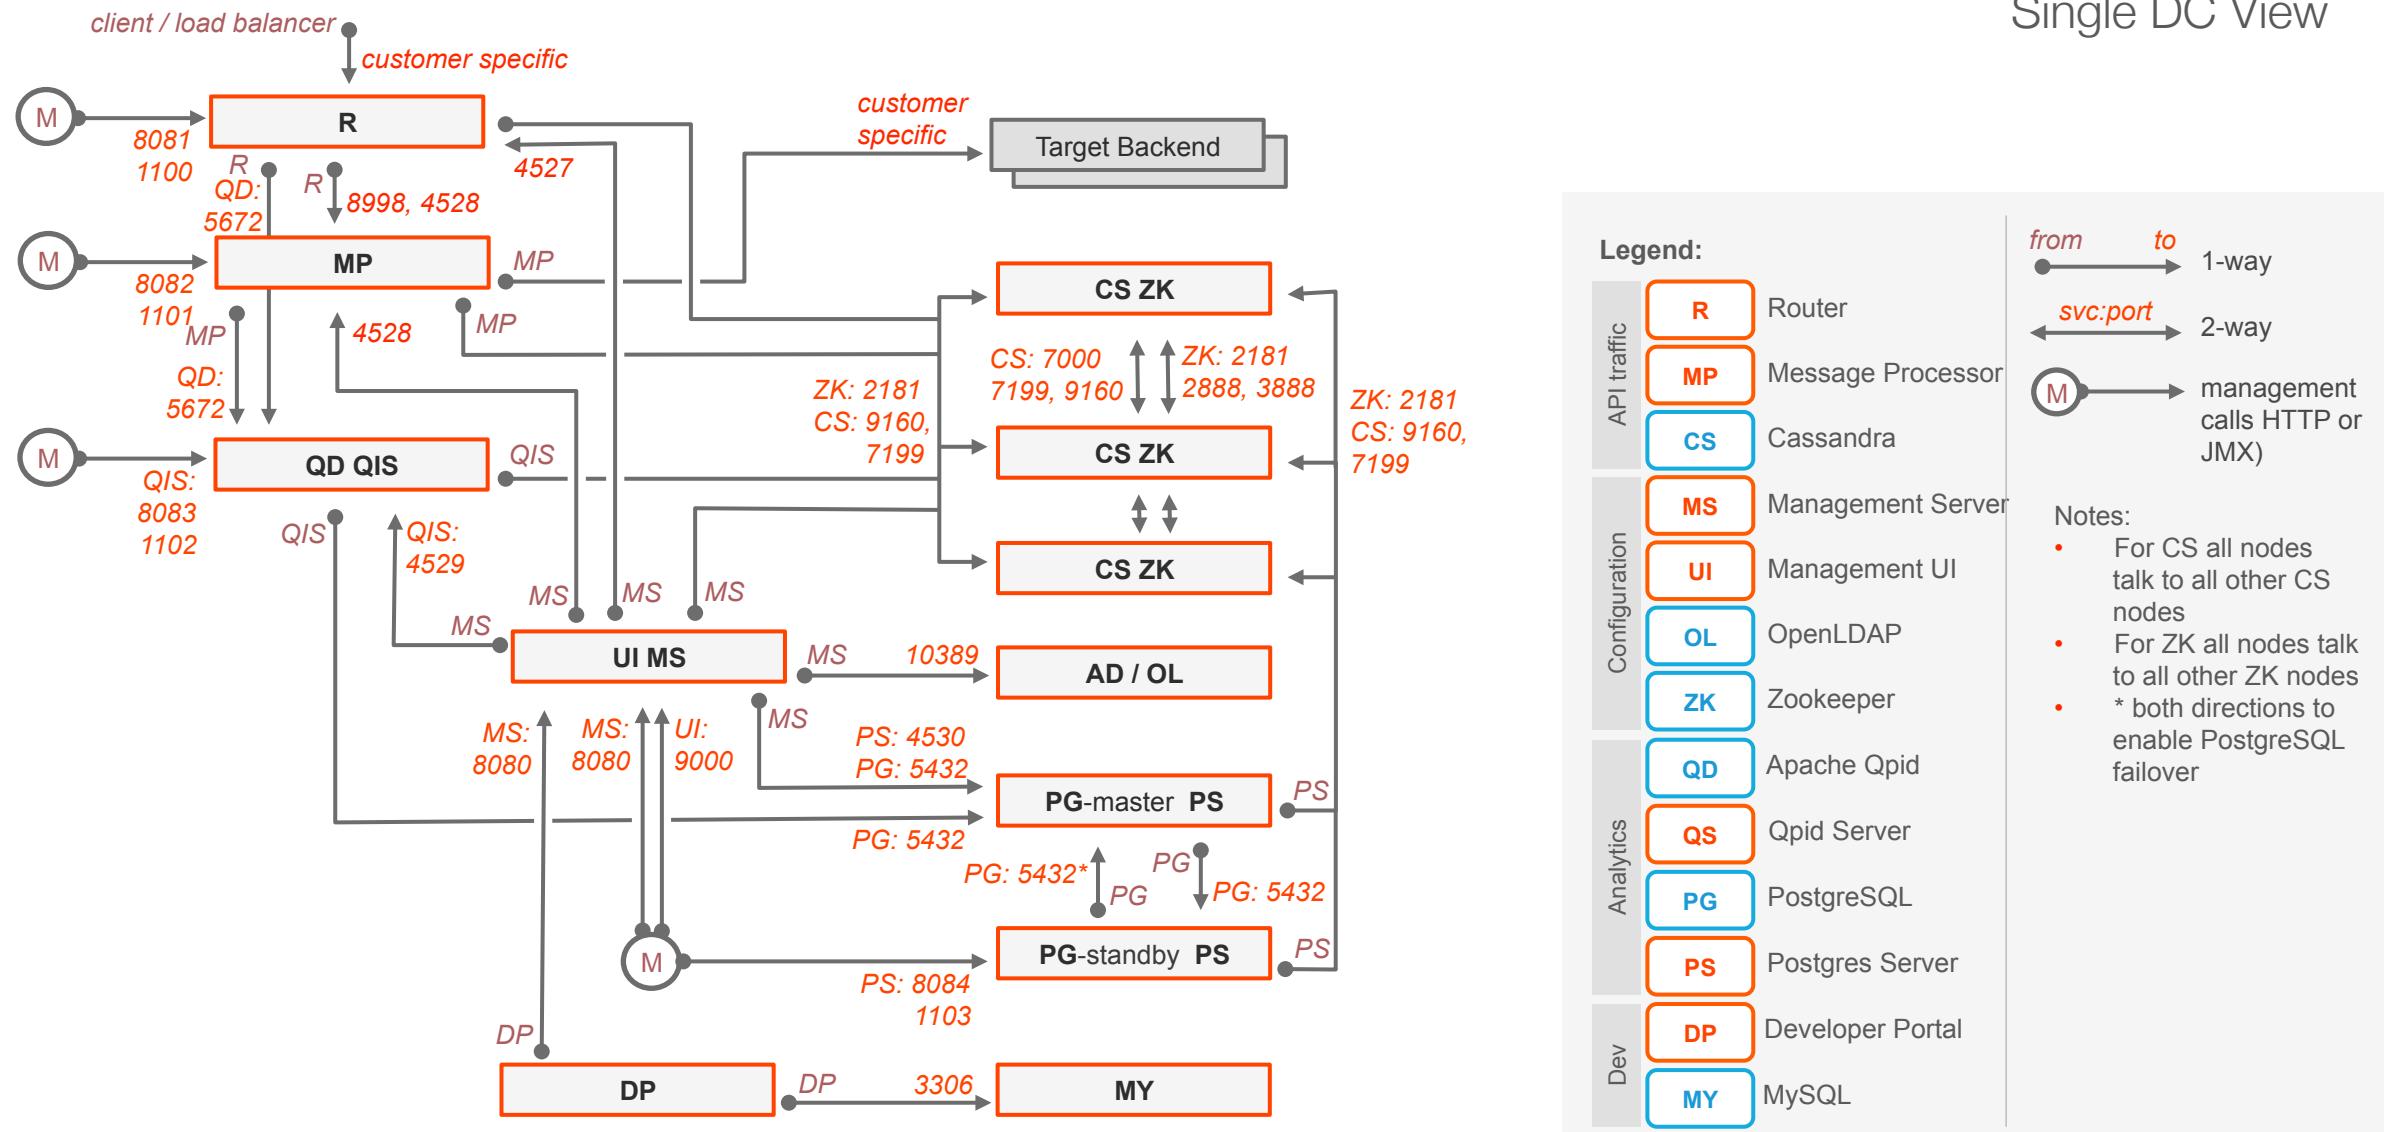
- Analytics data is partitioned per Environment.
- Analytics data size generated by MPs per transaction is about 1kb.
- PostgreSQL contains Raw Data Tables and Aggregated Data Tables. Raw Data Tables grow as analytics data is collected. Appropriate data retention policy and purge processes are required.
- Custom reports allow customers to define user-defined queries which run against raw data.

Apigee Edge Architecture – Network Zoning



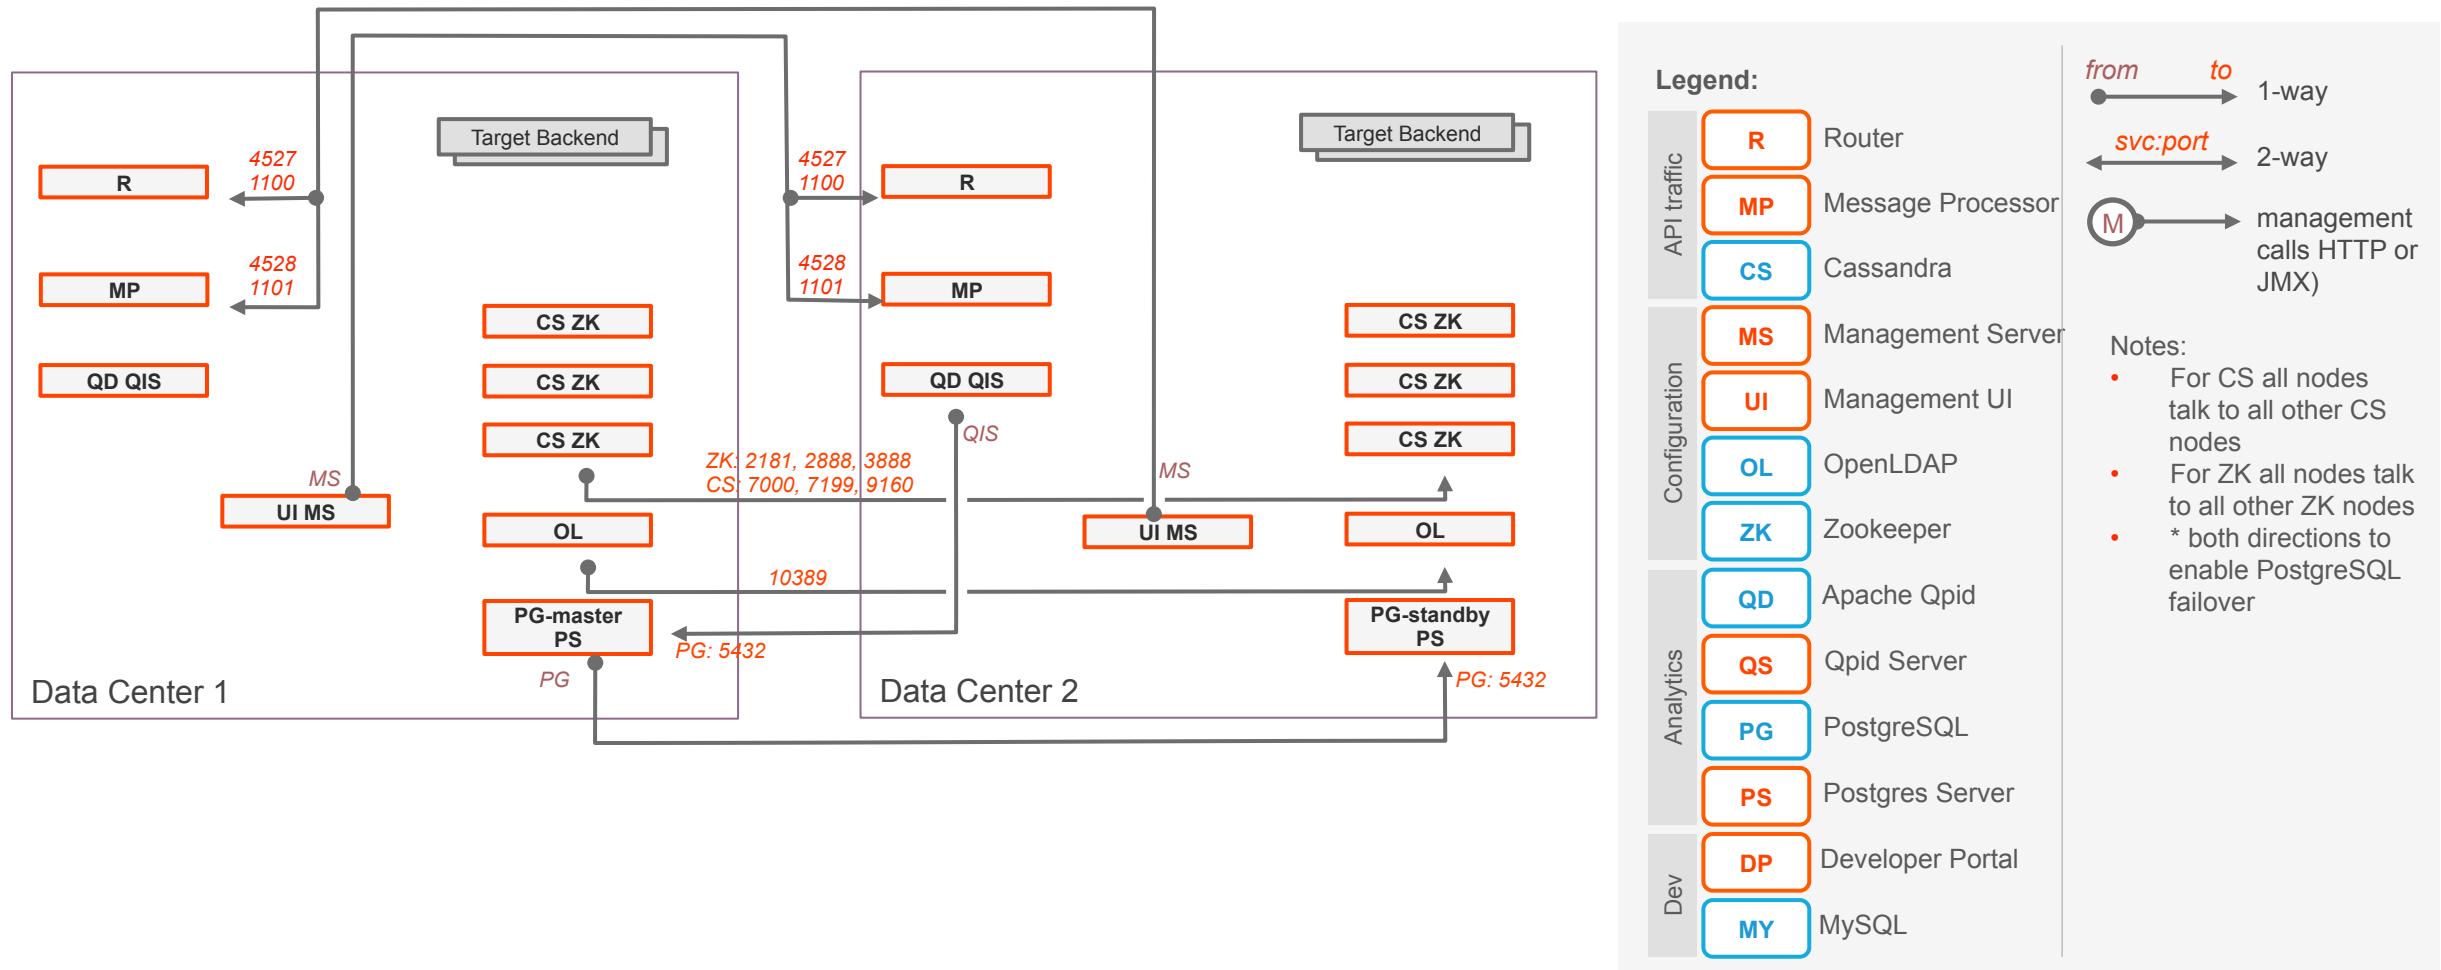
- Edge architecture does not impose network zone requirements. Network zoning will be driven by customer operation and security requirements.
- Firewalls and security appliances between zones should accommodate the connectivity and traffic characteristics of Apigee components without adding latency overhead.
- Keep API traffic as direct as possible. Keep MP dependencies close.
- Pay attention to cross-zone connections.

Apigee Edge Architecture – Components Connectivity



Apigee Edge Architecture – Components Connectivity

Multi DC View





60 mins break

Agenda

09:30 – 10:00

Apigee Overview

10:00 – 11:00

Apigee Private Cloud Architecture

BREAK (15 MINS)

11:15 – 12:00

Apigee Private Cloud Architecture (cont)

BREAK – LUNCH (60 MINS)

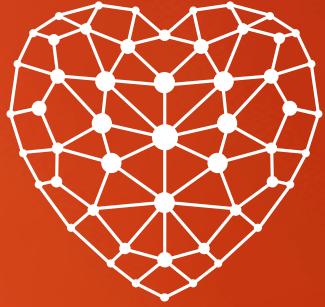
13:00 – 15:45

Apigee Edge Install

BREAK (15 MINS)

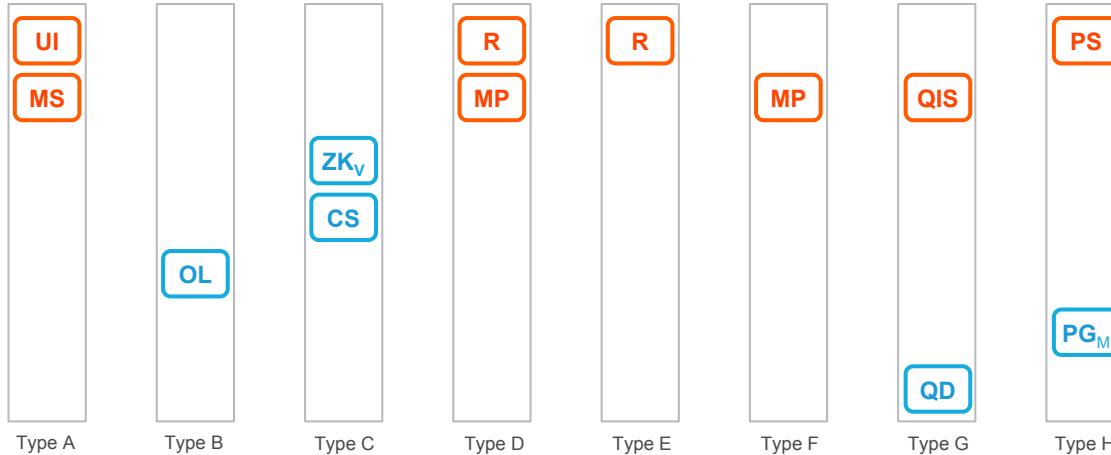
16:00 – 17:30

Platform Operations



Apigee Edge Install

System Requirements – Hardware Specification



- Cassandra
 - Local storage with SSD or fast HDD supporting 2000 IOPS
- PostgreSQL
 - Local storage with SSD or fast HDD supporting 1000 – 8000 IOPS
 - CPU, Memory and Size the storage according to the data retention requirements. Implement archive / purge process and keep less data for best performance.
 - Operations Guide provides formula to estimate storage needs.

Type	Components	CPU	Memory	Disk
A	UI, MS	2 core	4 GB	60 GB
B	OL	2 core	4 GB	60 GB
C	ZK, CS	8 core	16 GB	250 GB local storage with SSD or fast HDD supporting 2000 IOPS
D	R, MP	4/8 core	8/16 GB	100 GB
E	R	4 core	8 GB	60 GB
F	MP	8 core	16 GB	100 GB
G	QIS, QD	8 core	16 GB	500 GB local storage with SSD or fast HDD supporting 1000 IOPS
H	PS, PG	8 core	16 GB	500 GB to1 TB local storage with SSD or fast HDD supporting 4000-8000 IOPS

Legend:	R Router	MS Management Server	DP Developer Portal	MY MySQL	OL Openldap	SV Server/Virtual Machine
	MP Message Processor	PS Postgres Server	BA BaaS Server	ZK Zookeeper	PG PostgreSQL	POD POD
	UI Enterprise UI	QIS Qpid/Ingest Server		CS Cassandra	QD Apache Qpid	

System Requirements – Software Dependencies

Operating System

- Red Hat Enterprise Linux (64-bit):
 - 6.3, 6.4, 6.5, 6.6, 7.0
- CentOS (64-bit):
 - 6.3, 6.4, 6.5, 6.6, 7.0
- Oracle Linux (64-bit):
 - 6.5

JDK

- Oracle JDK 1.7
- OpenJDK 7

SSL/TLS

- 1.0
- 1.2

Other software

- Cassandra 2.0.15
- Zookeeper 3.4.5
- QPID 0.14
- PostgreSQL 9.3
- Play (UI) 2.3.4
- OpenLDAP 2.4

<https://apigee.com/docs/api-services/reference/supported-software>

Private Cloud Installer – Software Dependencies

- awk
- basename
- bash
- chkconfig
- curl
- date
- dirname
- echo
- expr
- grep
- hostname
- id
- ls
- perl
- pgrep (from procps)
- ps
- pwd
- python
- rpm
- rpm2cpio
- sed
- sudo
- tar
- tr
- uname
- unzip
- useradd
- wc
- yum

Complete list of prerequisites can be found in Apigee Edge Install and Configuration Guide, Page 19.

In addition to the tools above, some nodes require the installation of additional software components such as:

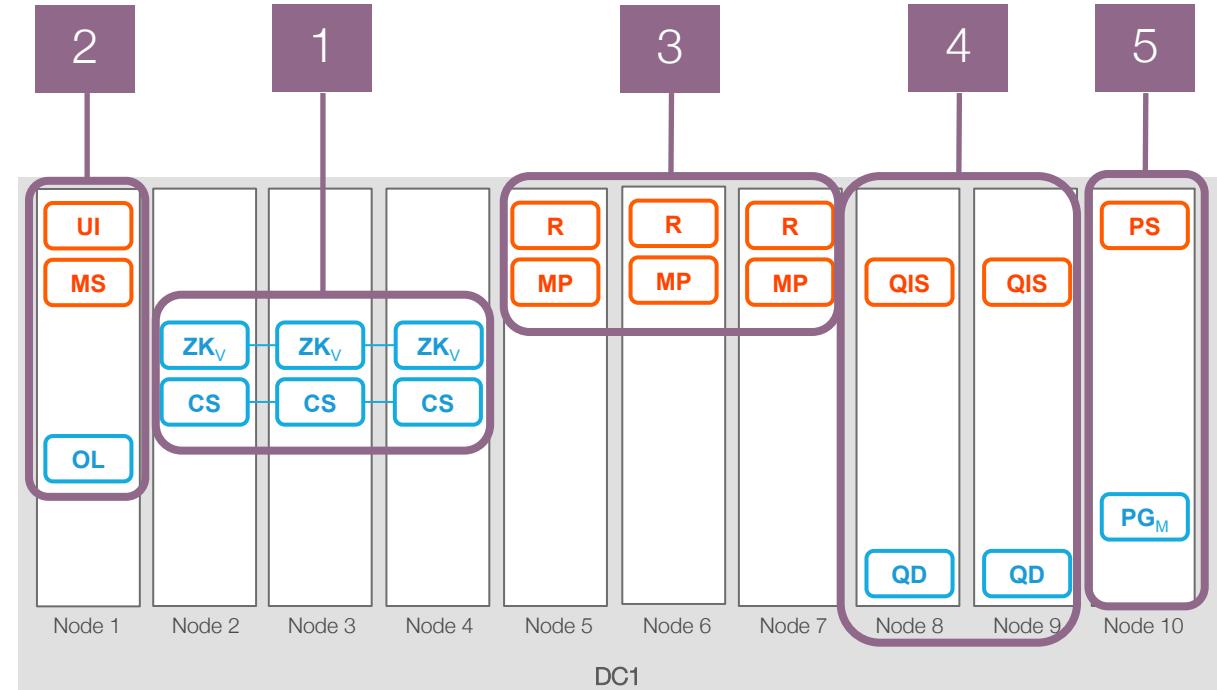
- ntp (all nodes)
- rsync (PostgreSQL nodes)
- openldap-clients openldap-servers (Management Server node)

Apigee Edge Installation Process

1. Install data store hosts
2. Install management hosts
 - Install LDAP first if using standalone LDAP hosts
3. Install router, message processor
4. Qpid hosts
5. PostgreSQL hosts
6. Configure PostgreSQL replication, if needed
7. Create organization(s) and environment(s)

If applicable, install:

7. Developer portal
8. Install monetization
9. App services (BaaS)



Legend:	R Router	MS Management Server	DP Developer Portal	MY MySQL	OL Openldap	Server/Virtual Machine
	MP Message Processor	PS Postgres Server	BA BaaS Server	ZK Zookeeper	PG PostgreSQL	POD
	UI Enterprise UI	QIS Qpid/Ingest Server		CS Cassandra	QD Apache Qpid	

Private Cloud – Software Install

Apigee Private Cloud (apigee-edge-4.15.07.00.zip) has everything required to install and configure Apigee Edge and BaaS components except for the Developer Portal. The Developer Portal is distributed using a different package (DeveloperServices-4.15.07.00.tar).

Apigee Edge installation steps:

1. Acquire a license key and copy it to all nodes.
2. Download the Apigee Private Cloud from <ftp.apigee.com> and copy it to all nodes.
3. Installation downloads and installs required system software via Yum. It requires Internet connection or local repository.
4. Unzip apigee-edge-4.<YY>.<MM>.<V>.zip
5. Run the primary installation script:

```
/<unzip-location>/apigee-edge-4.<YY>.<MM>.<V>/apigee-install.sh -j /usr/java/default  
-r <inst-root> -d <data-root>
```

Private Cloud – Software Setup

1. Once Apigee binaries are installed, run:

```
/<inst-root>/apigee4/share/installer/apigee-setup.sh
```

2. Setup script requires a profile type and information about the system, provided interactively or via response file (silent install).
3. Create an organization, environment(s) and organization administrator, run:

```
/<inst-root>/apigee4/bin/setup-org.sh
```

Private Cloud – Silent Install

- Silent install provides a way for specifying, in advance, all necessary values used by apigee-setup.sh
- Silent install provides a way for specifying, in advance, all necessary values used by apigee-setup.sh

```
<inst-root>/apigee4/share/installer/  
apigee-setup.sh -p ds -f <response-  
file-name>
```

- The response files contain a number of variables definitions to be used by apigee-setup.sh. A response file per DC/Region will be needed since some values are unique per DC/Region.

Single machine setups:

- aio = All In One (Gateway and Analytics Standalone)

Cluster node setup for ZooKeeper and Cassandra (min 3 nodes):

- ds = Datastore Cluster Node

LDAP setup for OpenLDAP:

- ld = LDAP Node

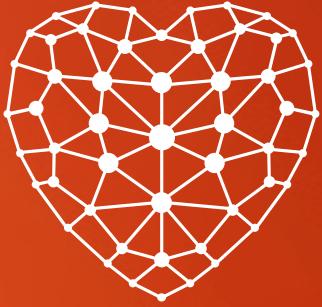
Separate components setup:

- ld = LDAP Node
- ms = Gateway Management Server
- r = Gateway Router
- mp = Gateway Message Processor
- rmp = Gateway Router and Message Processor
- qs = Analytics Qpid Server
- ps = Analytics Postgres Server
- mo = Monetization Server

Note: The list above doesn't show all the available options. Refers to Apigee Edge Install and Configuration Guide, section Basic Host Installation, page 35 for details.

Private Cloud – Silent Response File

- IP1=
- IP2=
- IP3=
- IP4=
- IP5=
- HOSTIP=\$(hostname -i)
- MSIP=\$IP1
- ADMIN_EMAIL=
- APIGEE_ADMINPW=
- LICENSE_FILE=/root/opdk/license.txt
- USE_LDAP_REMOTE_HOST=n
- LDAP_TYPE=1
- APIGEE_LDAPPW=
- ENABLE_AX=y
- MP POD=gateway
- REGION=dc-1
- USE_ZK_CLUSTER=y
- ZK_HOSTS="\$IP1 \$IP2 \$IP3"
- ZK_CLIENT_HOSTS="\$IP1 \$IP2 \$IP3"
- USE_CASS_CLUSTER=y
- CASS_HOSTS="\$IP1:1,1 \$IP2:1,1 \$IP3:1,1"
- CASS_USERNAME=
- CASS_PASSWORD=
- SKIP_SMTP=y
- SMTPHOST=
- SMTPPORT=25
- SMTPUSER=
- SMTTPASSWORD=
- SMTPSSL=n
- BIND_ON_ALL_INTERFACES=y



Apigee Edge Install Lab

Apigee Edge Install Lab

Scope

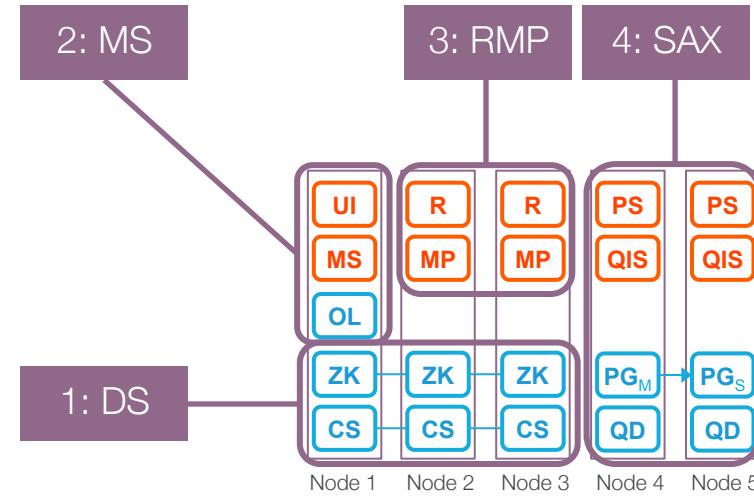
- 5 nodes Apigee Edge Install
- VMs instances running in Amazon AWS
- Installation and setup using silent install process.
- Setup order DS, MS, RMP, SAX

AWS Instances IP

- <http://iloveapis.com/workshops/ops/>

Documentation

- ApigeeEdgePrivateCloud-Install-Config-Guide.pdf
 - Page 35 - Basic Host Installation
 - Page 47 – Onboarding
 - Page 120 - Appendix A: Silent Installation
 - Page 121 - 5-host Clustered Installation



Apigee Edge Install Lab - Prerequisites

Preparation

- A. Software, sample install file and JDK are located in /root/opdk/
- B. Install prerequisites
 - o In all nodes:
 - Install JDK (rmp -ivh <rmp>)
 - yum -y install ntp
 - o In node 1:
 - yum -y install openldap-clients openldap-servers
 - o In node 4 and 5:
 - yum -y install rsync
- C. In all nodes, `unzip apigee-edge-4.15.07.00.zip`
- D. Update silent install file with [private IP addresses](#) corresponding to your AWS instances. Copy silent install files to all nodes.

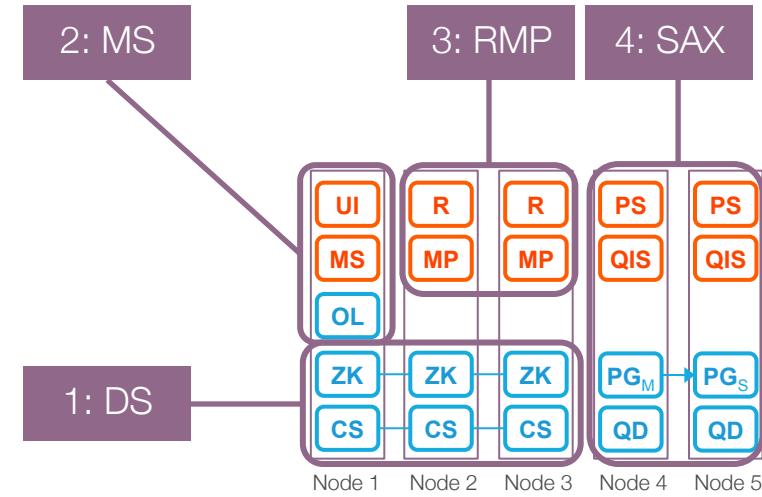
Additional considerations:

- For the purpose of the lab, we took care of some prerequisites. When planning for real world installation, prerequisites described in Apigee Edge Install and Configuration Guide, Page 19-35 should be completed.

Apigee Edge Install Lab – Installation and Setup

Install Steps

- A. Unzip apigee-edge-4.15.07.00.zip
- B. `./apigee-install.sh -j /usr/java/default -r /opt -d /opt`



Setup Steps

1. `/opt/apigee4/share/installer/apigee-setup.sh -p ds -f /root/opdk/response.txt`
2. `/opt/apigee4/share/installer/apigee-setup.sh -p ms -f /root/opdk/response.txt`
3. `/opt/apigee4/share/installer/apigee-setup.sh -p rmp -f /root/opdk/response.txt`
4. `/opt/apigee4/share/installer/apigee-setup.sh -p sax -f /root/opdk/response.txt`

Apigee Edge Install Lab – Installation and Setup

Setup Master-Standby Replication for PostgreSQL

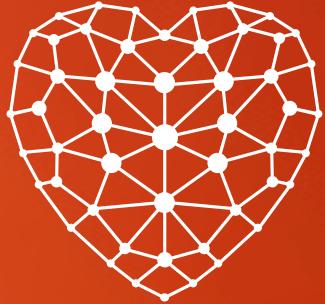
- Apigee Edge Install and Configuration Guide, Page 59

Onboarding

- Apigee Edge Install and Configuration Guide, Page 47
- /opt/apigee4/bin/setup-org.sh

Access Apigee Edge

- <http://<management-server-public-ip>:9000>



Platform Operations

Agenda

09:30 – 10:00

Apigee Overview

10:00 – 11:00

Apigee Private Cloud Architecture

BREAK (15 MINS)

11:15 – 12:00

Apigee Private Cloud Architecture (cont)

BREAK – LUNCH (60 MINS)

13:00 – 15:45

Apigee Edge Install Lab

BREAK (15 MINS)

16:00 – 17:30

Platform Operations

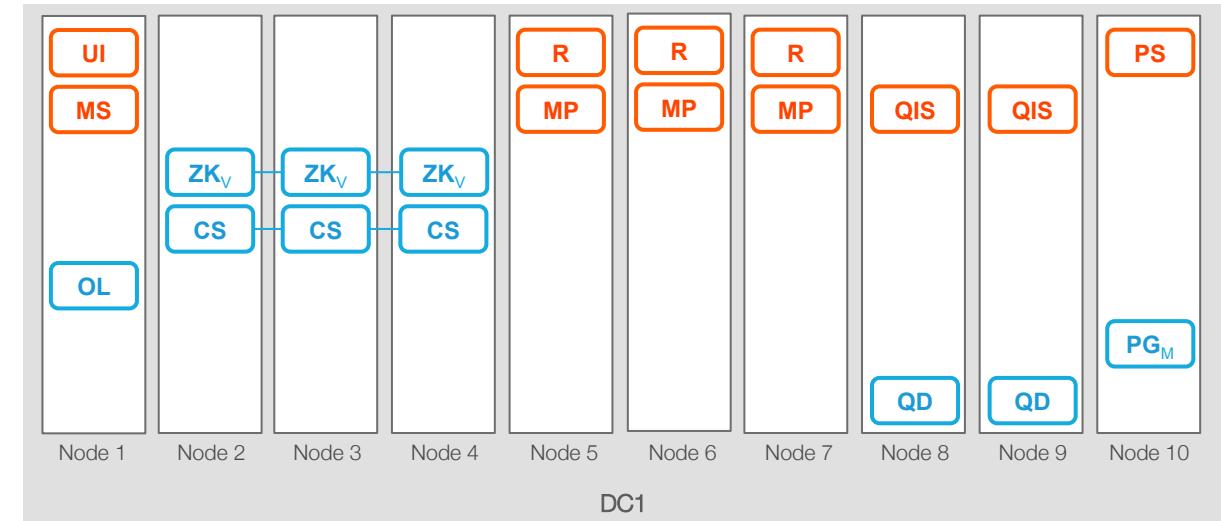
Platform Operations - Walkthrough

- Directory Structure
 - /<inst-root>/apigee4
- Start / Stop / Status
 - /<inst-root>/apigee4/bin/all-start.sh
 - /<inst-root>/apigee4/bin/all-status.sh
 - /<inst-root>/apigee4/bin/all-stop.sh
- Logs file
 - /<inst-root>/apigee4/var/log
- Management UI
 - <http://<management-server-host>:9000>
- Management API
 - <http://apigee.com/docs/management/apis>
 - <http://<management-server-host>:8080>



Platform Operations – Backup and Restore

- Backup/restore
 - Edge includes backup scripts for:
 - LDAP
 - Cassandra
 - ZooKeeper
 - PostgreSQL
 - UI (custom reports)
 - Component UUIDs
 - <inst-root>/apigee4/bin/backup.sh
- High Availability and Disaster Recovery
 - Multiple copies of the data
 - Eventual consistency
 - Horizontal scalability and resiliency



Platform Operations – Recurring Maintenance

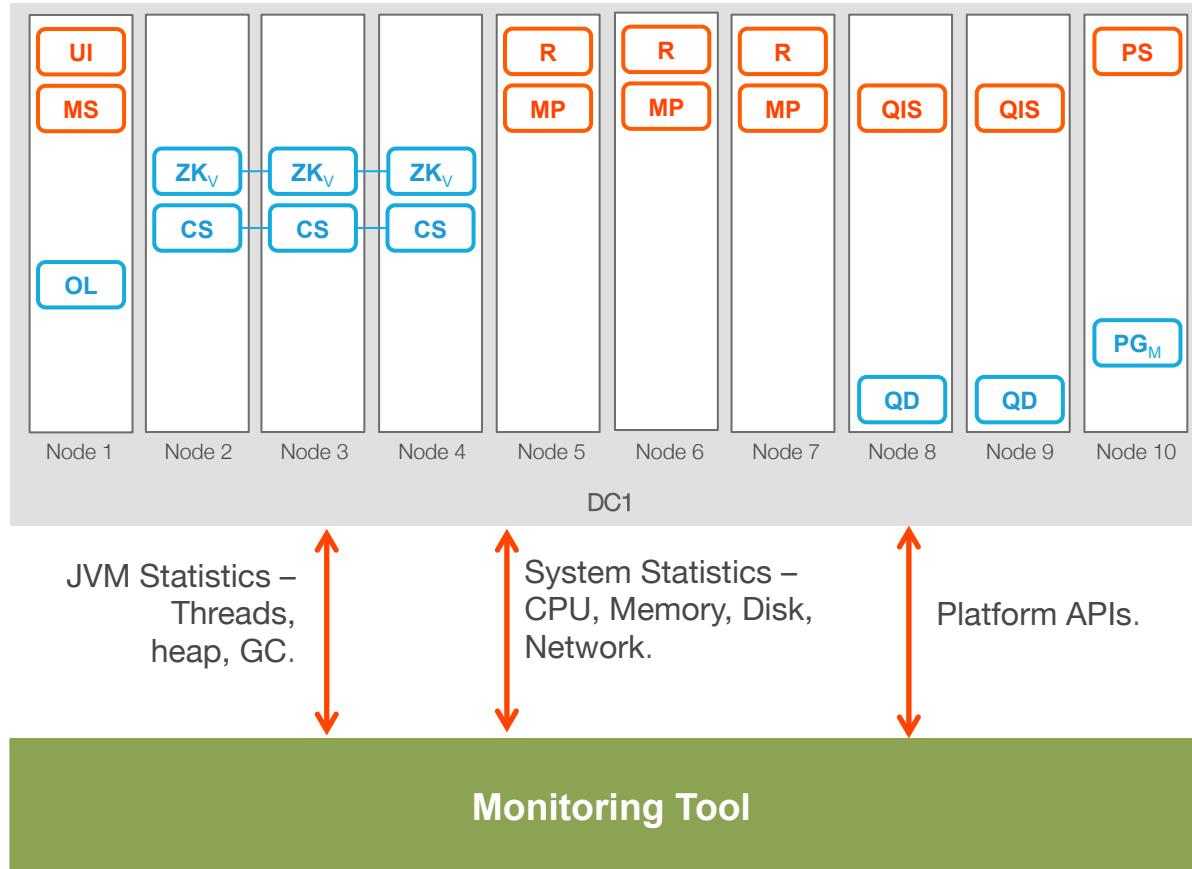
- Cassandra nodes repair

```
/opt/apigee4/share/apache-cassandra/bin/nodetool repair -pr
```

- Periodically prune detailed analytics data based on retention requirements

```
/opt/apigee4/bin/pg-data-purge.sh
```

Platform Operations – Component Level Monitoring



- System-Level Checks
 - CPU, Memory, Disk, Load, Network
- Process/Application Checks
 - Thread statistics, Memory utilization
- Components monitoring
 - JMX
 - Metrics
 - Management API (<http://<host>:<port>/v1/servers/self/up>)
 - Router health check of the Message Processor
 - Log monitoring
 - HeartBeat
 - Mark down/mark up events
 - Logging Policy

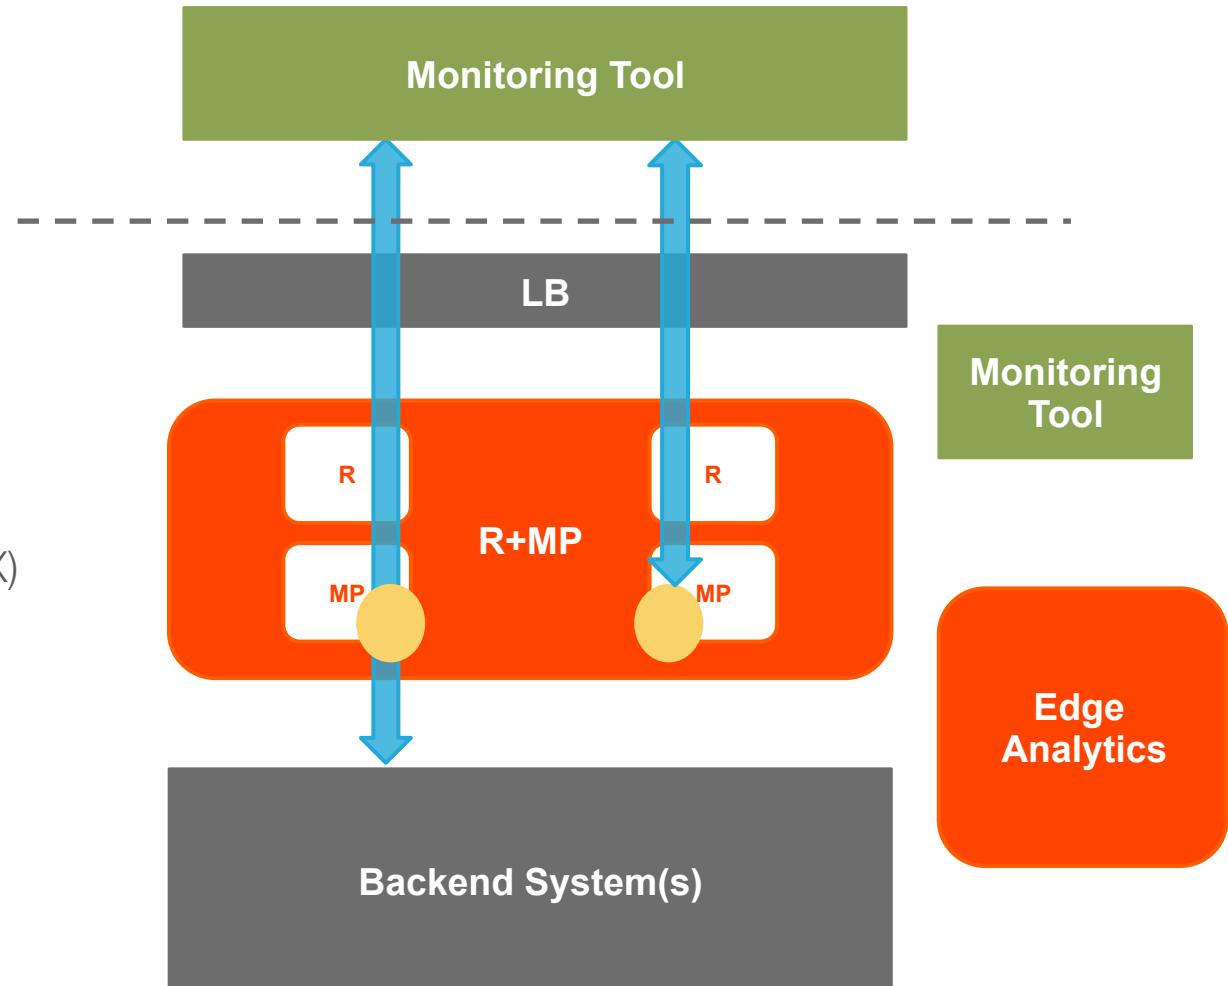
Legend:

R	Router
MP	Message Processor
UI	Enterprise UI
MS	Management Server
PS	Postgres Server
QIS	Qpid/Ingest Server
DP	Developer Portal
BA	BaaS Server
ZK	Zookeeper
CS	Cassandra
OL	Openldap
PG	PostgreSQL
QD	Apache Qpid

MY	MySQL
PG_M	PostgreSQL
QD	Apache Qpid
POD	POD

Platform Operations – APIs Monitoring

- API-Level Checks
 - API calls 20X responses
 - API call error within Apigee
 - API call error on backend
 - JMX
 - Metrics
- Message Flow Checks
 - Number of active clients
 - Number of responses (10X, 20X, 30X, 40X and 50X)
 - Connect failures
- API Health - <https://pages.apigee.com/api-health>



Legend:

R Router

MP Message Processor

UI Enterprise UI

MS Management Server

PS Postgres Server

QIS Qpid/Ingest Server

DP Developer Portal

BA BaaS Server

CS Cassandra

MY MySQL

ZK Zookeeper

QL Apache Qpid

OL Openldap

PG PostgreSQL

QD Apache Qpid

VM Server/Virtual Machine

POD POD

Takeaway

- Public Cloud = Private Cloud
- Architecture purpose built and proven
- Horizontally scalable, components can be scaled independently of each other
- Multi-tenant by design

Thank you!